

Units  
oooooo

Organization  
oooooooo

CPU  
oooooooo

RAM  
oooooo

Buses  
oooooooo

VMs  
oooooooooooo

# A Quick Primer on Computer Architecture

ComS 252 — Iowa State University

Andrew Miner

Units  
●ooooo

Organization  
oooooooo

CPU  
oooooooo

RAM  
oooooo

Buses  
oooooooo

VMs  
oooooooooooo

# SI multiplier prefixes — small

- ▶ SI: International System of Units

Units  
●ooooo

Organization  
oooooooo

CPU  
ooooooo

RAM  
oooooo

Buses  
ooooooo

VMs  
oooooooooooo

## SI multiplier prefixes — small

- ▶ SI: International System of Units
  - ▶ Shouldn't it be IS? or ISU?

## SI multiplier prefixes — small

- ▶ SI: International System of Units
    - ▶ Shouldn't it be IS? or ISU?
    - ▶ SI abbreviation is French: *Système International d'unités*
    - ▶ Basically, "the metric system"

# SI multiplier prefixes — small

- ▶ SI: International System of Units
  - ▶ Shouldn't it be IS? or ISU?
  - ▶ SI abbreviation is French: *Système International d'unités*
  - ▶ Basically, “the metric system”
- ▶ Important “small” prefixes in computing:

Prefix	Sym.	Multiplier
milli	m	$10^{-3}$
micro	$\mu$	$10^{-6}$
nano	n	$10^{-9}$
pico	p	$10^{-12}$

- ▶ Note: all symbols are small letters
- ▶ Typical for times, physical sizes
  - ▶ 45 nm CPU (45 nanometers: refers to “feature” size on CPU)
  - ▶ 15  $\mu$ s to send a message

### SI multiplier prefixes — large

- ▶ Important “large” prefixes in computing:

<b>Prefix</b>	<b>Sym.</b>	<b>Multiplier</b>
kilo	k	$10^3$
mega	M	$10^6$
giga	G	$10^9$
tera	T	$10^{12}$
peta	P	$10^{15}$
exa	E	$10^{18}$
zetta	Z	$10^{21}$

- ▶ Note: **most** symbols are capital letters
  - ▶ These are used for capacities, and frequencies
    - ▶ 1 Tb drive (1 Terabyte: 1 trillion bytes)
    - ▶ 2 Ghz clock (2 Gigahertz: 2 billion cycles per second)

# Binary notation

- ▶ Humans write numbers in *decimal* notation:
  - ▶ Ten possibilities for each digit (0, 1, ..., 8, 9)
  - ▶  $507 = 5 \cdot 10^2 + 0 \cdot 10^1 + 7 \cdot 10^0$
  - ▶ How many different numbers can  $n$  digits represent?

# Binary notation

- ▶ Humans write numbers in *decimal* notation:
  - ▶ Ten possibilities for each digit (0, 1, ..., 8, 9)
  - ▶  $507 = 5 \cdot 10^2 + 0 \cdot 10^1 + 7 \cdot 10^0$
  - ▶ How many different numbers can  $n$  digits represent?  $10^n$

# Binary notation

- ▶ Humans write numbers in *decimal* notation:
  - ▶ Ten possibilities for each digit (0, 1, ..., 8, 9)
  - ▶  $507 = 5 \cdot 10^2 + 0 \cdot 10^1 + 7 \cdot 10^0$
  - ▶ How many different numbers can  $n$  digits represent?  $10^n$
- ▶ No fundamental reason that the “base” needs to be 10
- ▶ Computers represent numbers in *binary* notation:
  - ▶ Two possibilities for each digit (0, 1)
  - ▶  $507_{10} = 1 \cdot 2^8 + 1 \cdot 2^7 + 1 \cdot 2^6 + 1 \cdot 2^5 + 1 \cdot 2^4 + 1 \cdot 2^3 + 0 \cdot 2^2 + 1 \cdot 2^1 + 1 \cdot 2^0$
  - ▶  $507_{10} = 111111011_2$
  - ▶  $n$  binary digits (bits) can represent

# Binary notation

- ▶ Humans write numbers in *decimal* notation:
  - ▶ Ten possibilities for each digit (0, 1, ..., 8, 9)
  - ▶  $507 = 5 \cdot 10^2 + 0 \cdot 10^1 + 7 \cdot 10^0$
  - ▶ How many different numbers can  $n$  digits represent?  $10^n$
- ▶ No fundamental reason that the “base” needs to be 10
- ▶ Computers represent numbers in *binary* notation:
  - ▶ Two possibilities for each digit (0, 1)
  - ▶  $507_{10} = 1 \cdot 2^8 + 1 \cdot 2^7 + 1 \cdot 2^6 + 1 \cdot 2^5 + 1 \cdot 2^4 + 1 \cdot 2^3 + 0 \cdot 2^2 + 1 \cdot 2^1 + 1 \cdot 2^0$
  - ▶  $507_{10} = 111111011_2$
  - ▶  $n$  binary digits (bits) can represent  $2^n$  different numbers

# Hexadecimal

- ▶ Base 16: digits are (0, 1, ..., 9, a, b, c, d, e, f)
- ▶ 1 hexadecimal digit equals four bits
- ▶ Easy to convert between hex and binary

$$0_{16} = 0000_2 \quad 1_{16} = 0001_2 \quad 2_{16} = 0010_2 \quad 3_{16} = 0011_2$$

$$4_{16} = 0100_2 \quad 5_{16} = 0101_2 \quad 6_{16} = 0110_2 \quad 7_{16} = 0111_2$$

$$8_{16} = 1000_2 \quad 9_{16} = 1001_2 \quad a_{16} = 1010_2 \quad b_{16} = 1011_2$$

$$c_{16} = 1100_2 \quad d_{16} = 1101_2 \quad e_{16} = 1110_2 \quad f_{16} = 1111_2$$

- ▶  $507_{10} = 1,1111,1011_2 = 1fb_{16}$
- ▶ Often written as 0x1fb or x1fb
- ▶ MAC addresses: 12 hex digits
  - ▶ Much easier to write than 48 bits

# Octal

- ▶ Base 8: digits are (0, 1, . . . , 7)
- ▶ 1 octal digit equals three bits — also easy to convert
  - $0_8 = 000_2$
  - $1_8 = 001_2$
  - $2_8 = 010_2$
  - $3_8 = 011_2$
  - $4_8 = 100_2$
  - $5_8 = 101_2$
  - $6_8 = 110_2$
  - $7_8 = 111_2$
- ▶  $507_{10} = 111,111,011_2 = 773_8$
- ▶ More obscure than hexadecimal, but still used

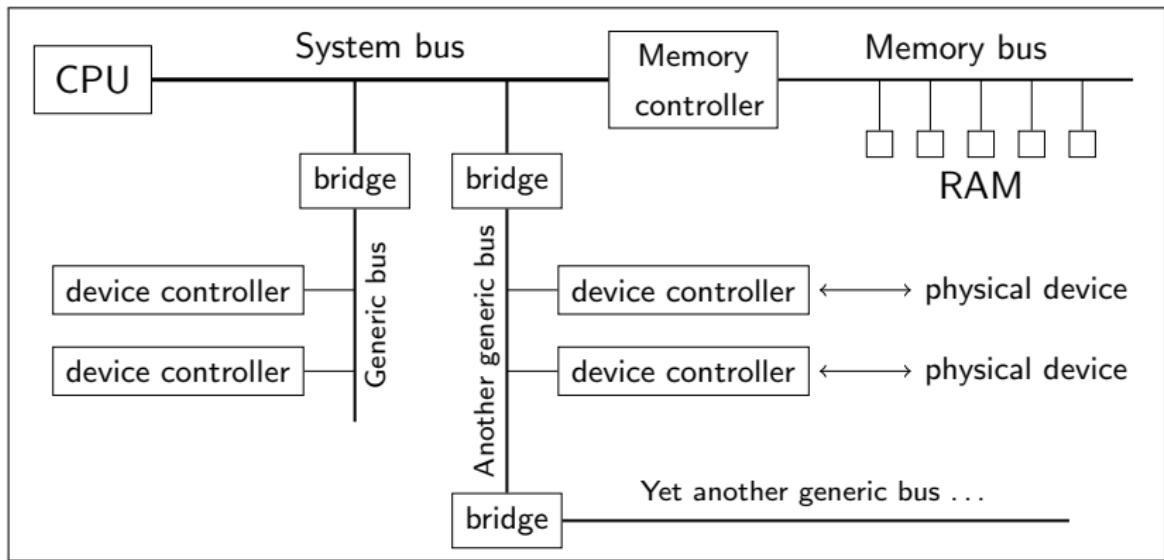
# Binary prefixes

- ▶  $10^n$  is a nice round number in decimal notation
- ▶  $2^n$  is a nice round number in binary notation
- ▶ In computing, it is often easier to group  $2^n$  things
  - ▶ E.g.,  $2^{10} = 1024$  bytes is more natural than 1000 bytes
- ▶ Slightly different prefixes used:

Confusing	Preferred	Sym.	Multiplier
kilo	Kibi	Ki	$2^{10} = 1024^1$
mega	Mebi	Mi	$2^{20} = 1024^2$
giga	Gibi	Gi	$2^{30} = 1024^3$
	Tebi	Ti	$2^{40} = 1024^4$

- ▶ Can lead to issues
  - ▶ Manufacturer says 1Tb drive (1,000,000,000,000 bytes)
  - ▶ Disk utility says 931 Gib drive:  $931 \times 2^{30} \approx 10^{12}$   
(even less after formatting)

# Simple, abstract view of a computer



Units  
oooooo

Organization  
o●oooooo

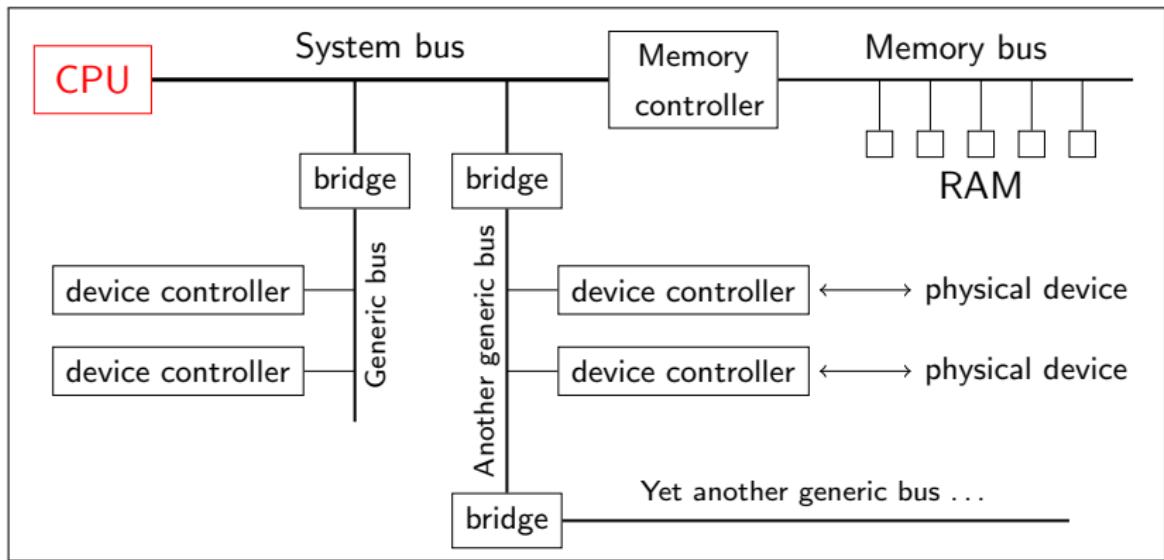
CPU  
oooooooo

RAM  
oooooo

Buses  
ooooooo

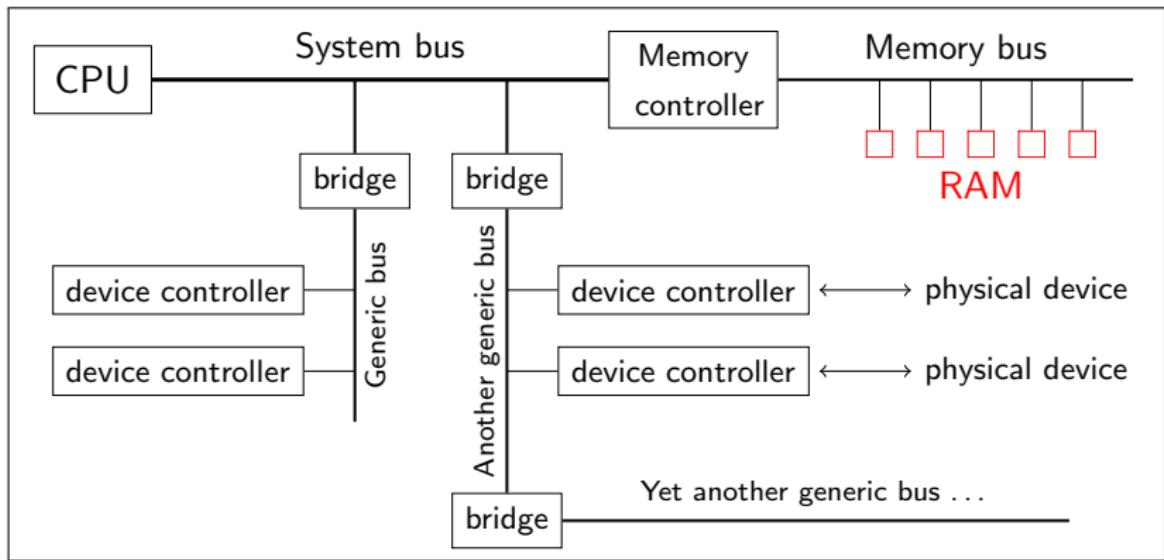
VMs  
oooooooooooo

# CPU: Central Processing Unit



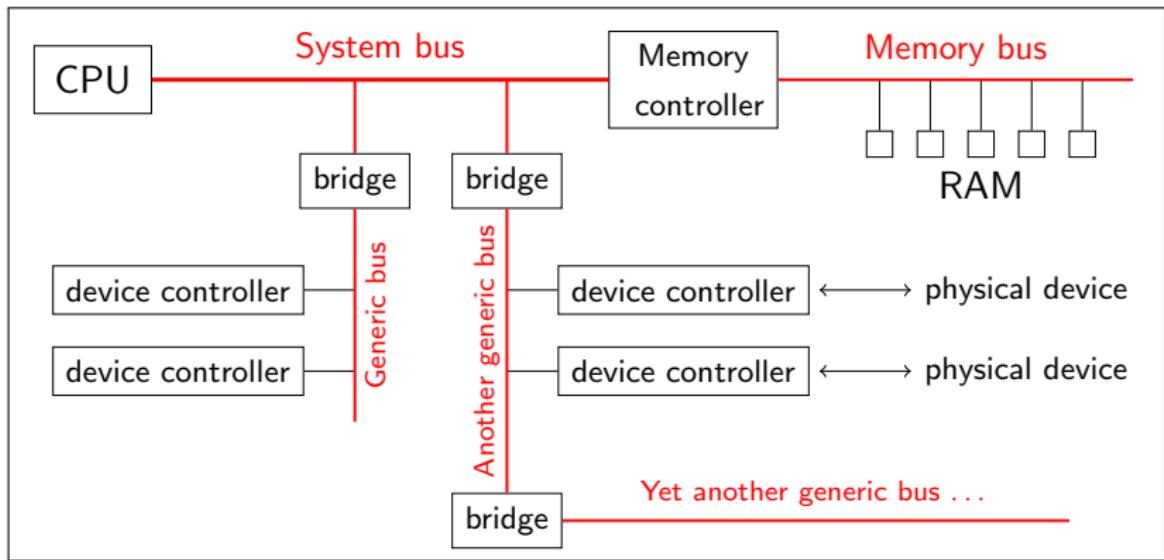
- ▶ Executes machine language instructions

# RAM: Random Access Memory



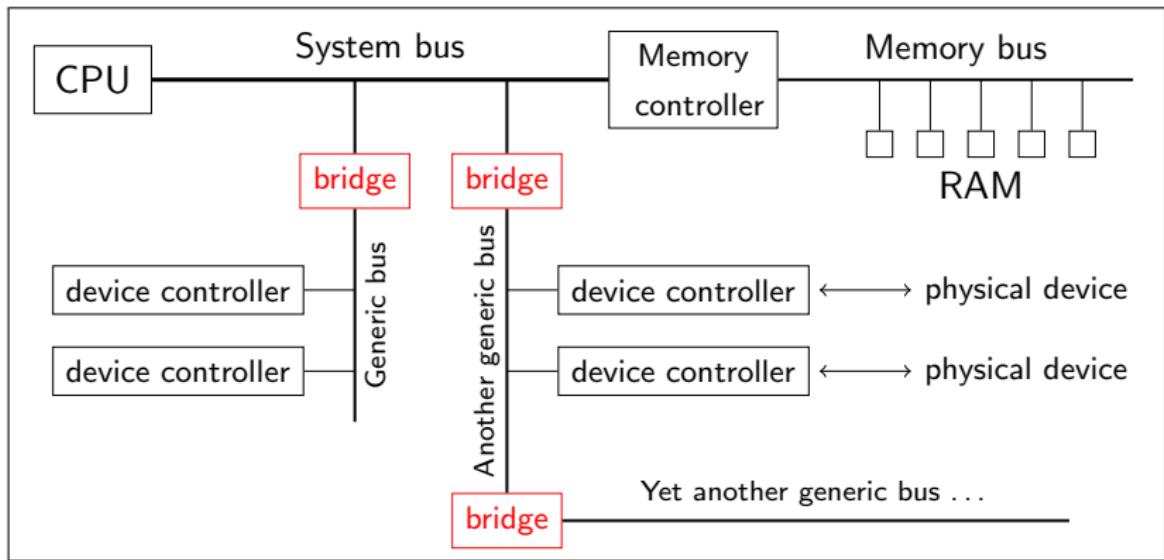
- ▶ Stores data
- ▶ Unit of storage: *byte* (8 bits)
- ▶ Each byte has an integer *address*

# Buses



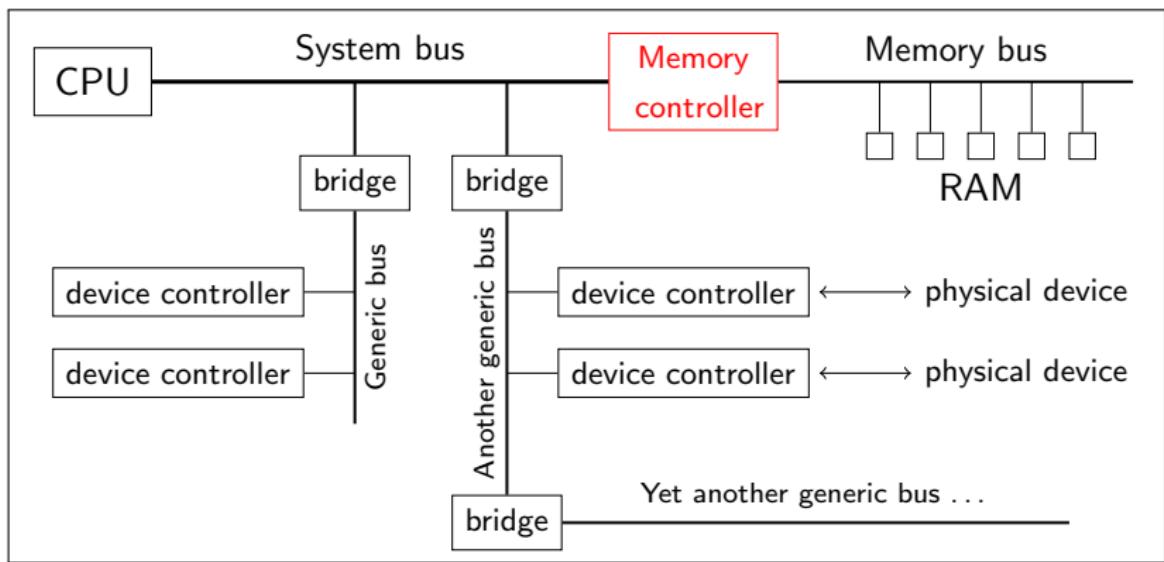
- ▶ Wires: allow data transfer
- ▶ Protocol: rules for transferring data

# Bridges



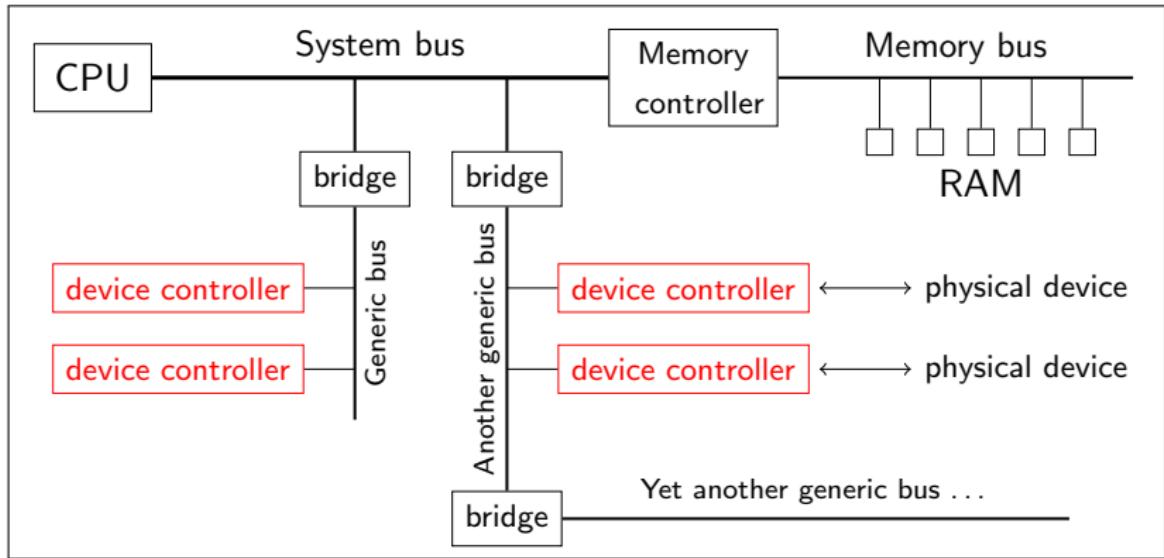
- ▶ Convert from one bus type to another

# Memory controller



- ▶ Specialized bridge for RAM
- ▶ Has additional functionality (e.g., for paging)
- ▶ Details are “beyond the scope of this class”

# Device controllers



- ▶ Connect to physical devices

# Instructions

CPUs execute instructions in *machine language* (ML)

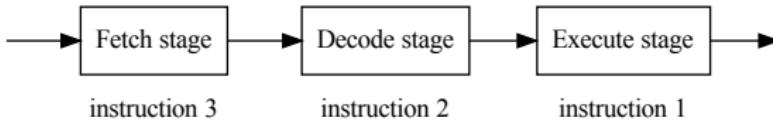
- ▶ Fairly simple, e.g.:
  - ▶ Copy integer at memory address 0x0c2b4938 into variable b
  - ▶ Subtract b from a
  - ▶ If a is zero, start executing instructions at address 0x003a247c
- ▶ High-level languages (like C) must be translated into ML  
(That's what a *compiler* does)
- ▶ **Instruction set**: set of instructions CPU can execute
- ▶ Family of CPUs: have backward compatible instruction sets
  - ▶ E.g., Intel 8086, Intel 80286, Intel 80386
- ▶ Different CPUs typically have different instruction sets
  - ▶ ML code for Motorola 68000 family  
**will not run** on Intel 8086 family

# Instruction Cycle

CPUs execute ML code in cycles, e.g.:

1. Fetch the next instruction to execute, from memory
2. Decode the instruction
  - ▶ Instructions are encoded into some number of bytes
  - ▶ Some instructions have “arguments”
    - ▶ E.g., copy into b from *where*?
3. Execute the instruction

Cycles are broken into **stages**, arranged in a **pipeline**:



Reality: pipelines are more complex than this

- ▶ RISC processors have 5 stage pipeline
- ▶ Pentium 3: 10 stage pipeline
- ▶ Pentium 4: 20 stage pipeline

# CPU clock

- ▶ CPUs have an internal **clock** (like a metronome)
  - ▶ Stages are synchronized
  - ▶ Work within stages is synchronized
  - ▶ To learn why, take CprE 281
- ▶ Completely different from **wall clock**
- ▶ Clock rate is measured in hertz
- ▶ Speed at which a CPU can execute a program depends on
  - ▶ Instruction set
  - ▶ Pipeline layout
  - ▶ Stage design
  - ▶ Other details we have not discussed yet...
  - ▶ Clock frequency
- ▶ Comparing clock frequencies is valid **only for similar CPUs**
  - ▶ 16 Mhz Intel 80386 **is slower than** 33 Mhz Intel 80386
  - ▶ 0.867 Ghz G4 **can run faster than** 2 Ghz Pentium 4

# Moore's Law

## Moore's Law

Predicts that the number of transistors that can be placed inexpensively on an integrated circuit doubles approximately every two years

- ▶ Not a *physical* law; rather, a “rule of thumb”
- ▶ Moore made the original prediction in 1965
  - ▶ Originally, doubled *every year*
  - ▶ Moore predicted it would hold for 10 years
  - ▶ In 1975, Moore altered it to *every two years*
  - ▶ Industry has treated Moore's law as a target
- ▶ Is the end near for Moore's law?
  - ▶ 2017: 5 nm chips announced (~ 30 billion transistors)
  - ▶ 2021: 2 nm chips announced (~ 50 billion transistors)
    - ▶ Human DNA strand is 2.5 nm in diameter

## Units

oooooo

## Organization



CPU  
0000●00

RAM  
00000

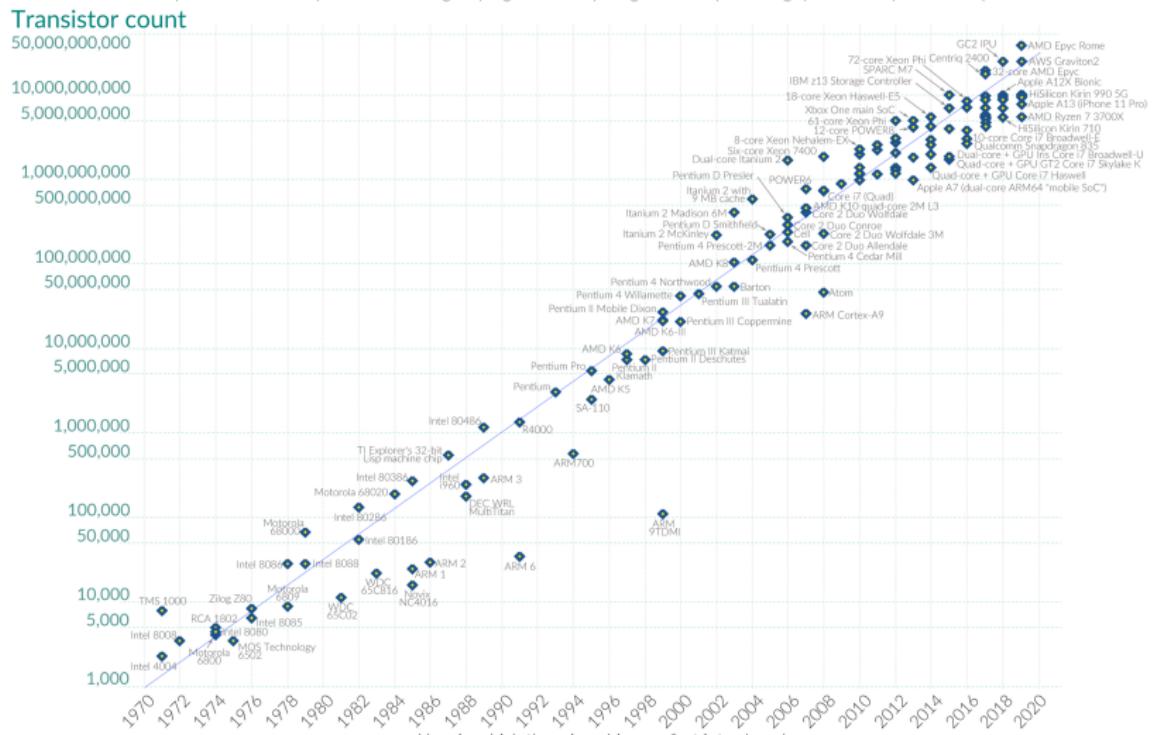
Buses  
ooooooo

VMs  
○○○○○○○○○○

Moore's Law: The number of transistors on microchips doubles every two years

Moore's law describes the empirical regularity that the number of transistors on integrated circuits doubles approximately every two years. This advancement is important for other aspects of technological progress in computing – such as processing speed or the price of computers.

Our World  
in Data



Data source: Wikipedia ([https://en.wikipedia.org/w/index.php?title=Transistor\\_count&oldid=1000000000](https://en.wikipedia.org/w/index.php?title=Transistor_count&oldid=1000000000))

[OurWorldInData.org](http://OurWorldInData.org) - Research and data to make progress against the world's largest problems.

Licensed under CC-BY by the authors Hannah Ritchie and Max Roser.

Units  
oooooo

Organization  
oooooooo

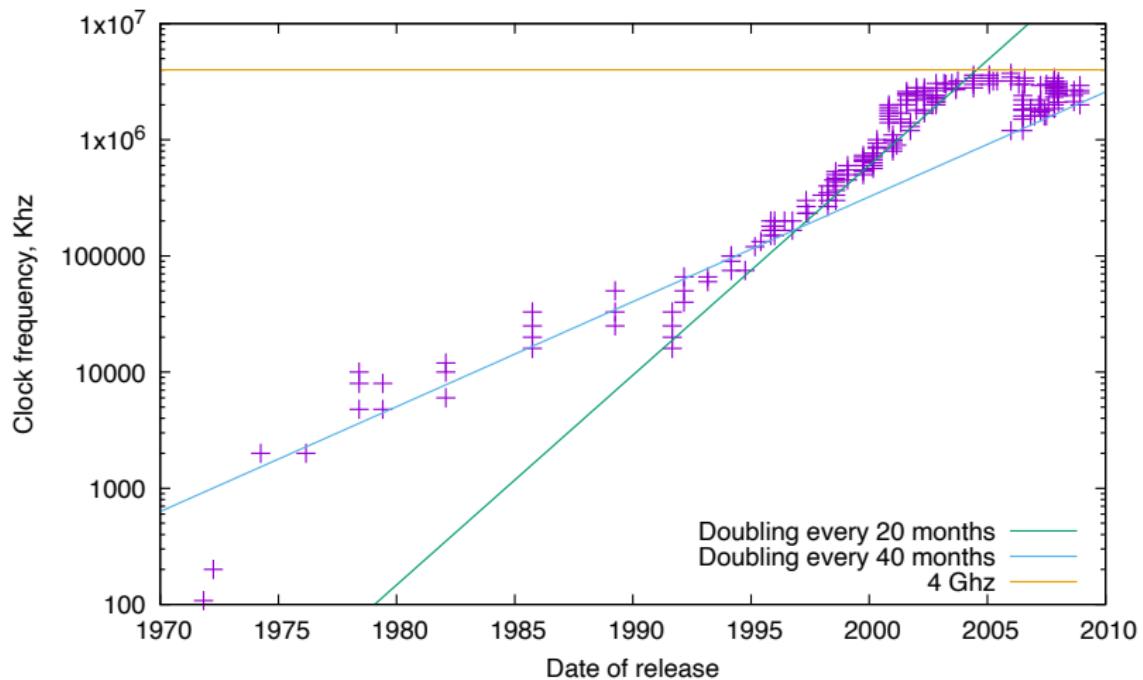
CPU  
ooooooo●o

RAM  
oooooo

Buses  
oooooooo

VMs  
oooooooooooo

# Clock frequencies in some Intel CPUs



# The shift to multi-core

- ▶ CPU clock frequencies have not increased (much) since 2005
  - ▶ Power / Heat is the main issue
- ▶ Manufacturers instead started increasing the number of *cores*
  - ▶ Effectively, each core is its own processor
- ▶ Huge gamble: nobody was asking for multi-core
- ▶ How do I keep 8 cores busy? What about 128 cores?
  - ▶ Image and video processing are easy to parallelize
  - ▶ Most people run several things at once
  - ▶ OS's will need smarter schedulers

# Random Access Memory

- ▶ Stores programs (ML instructions) and data
- ▶ Volatile: loses data when power is cut
- ▶ Each byte has a unique address
  - ▶ Consecutive integers from 0 to  $N - 1$
  - ▶ E.g., 1 Mb of RAM has addresses 0, 1, 2, ...,  $\text{fffff}_{16}$ 
    - ▶  $\text{fffff}_{16} = 1111\ 1111\ 1111\ 1111\ 1111_2 = 2^{20} - 1 = 1024^2 - 1$
- ▶ Memory is divided into *pages* of  $2^n$  bytes
  - ▶ Choice for  $n$  depends on processor architecture
  - ▶  $n = 12$  (typical) gives 4Kb pages
  - ▶ Using  $2^n$  means address can be split at bit boundary:

E.g., address 0x8e64a2: 1000 1110 0110 0100 1010 0010  
page number      offset within page

# 32-bit vs. 64-bit systems

- ▶ CPUs use *registers* to store integers
  - ▶ 32-bit CPU: (most) registers are 32 bits
  - ▶ 64-bit CPU: (most) registers are 64 bits
- ▶ So what?

# 32-bit vs. 64-bit systems

- ▶ CPUs use *registers* to store integers
  - ▶ 32-bit CPU: (most) registers are 32 bits
  - ▶ 64-bit CPU: (most) registers are 64 bits
- ▶ Address register size determines amount of addressable RAM

# 32-bit vs. 64-bit systems

- ▶ CPUs use *registers* to store integers
  - ▶ 32-bit CPU: (most) registers are 32 bits
  - ▶ 64-bit CPU: (most) registers are 64 bits
- ▶ Address register size determines amount of addressable RAM
  - ▶ 20 bits (old 8086):  $2^{20}$  bytes addressable = 1 Mb

# 32-bit vs. 64-bit systems

- ▶ CPUs use *registers* to store integers
  - ▶ 32-bit CPU: (most) registers are 32 bits
  - ▶ 64-bit CPU: (most) registers are 64 bits
- ▶ Address register size determines amount of addressable RAM
  - ▶ 20 bits (old 8086):  $2^{20}$  bytes addressable = 1 Mb
  - ▶ 32 bits:  $2^{32}$  bytes = 4 Gibi
    - ▶ **4 Gb limit** on 32-bit machines

# 32-bit vs. 64-bit systems

- ▶ CPUs use *registers* to store integers
  - ▶ 32-bit CPU: (most) registers are 32 bits
  - ▶ 64-bit CPU: (most) registers are 64 bits
- ▶ Address register size determines amount of addressable RAM
  - ▶ 20 bits (old 8086):  $2^{20}$  bytes addressable = 1 Mb
  - ▶ 32 bits:  $2^{32}$  bytes = 4 Gibi
    - ▶ **4 Gb limit** on 32-bit machines
  - ▶ 48 bits:  $2^{48}$  bytes = 256 Tebi

# 32-bit vs. 64-bit systems

- ▶ CPUs use *registers* to store integers
  - ▶ 32-bit CPU: (most) registers are 32 bits
  - ▶ 64-bit CPU: (most) registers are 64 bits
- ▶ Address register size determines amount of addressable RAM
  - ▶ 20 bits (old 8086):  $2^{20}$  bytes addressable = 1 Mb
  - ▶ 32 bits:  $2^{32}$  bytes = 4 Gibi
    - ▶ **4 Gb limit** on 32-bit machines
  - ▶ 48 bits:  $2^{48}$  bytes = 256 Tebi
  - ▶ 64 bits:  $2^{64}$  bytes = 16 Exabytes (16 billion Gb)

# 32-bit vs. 64-bit systems

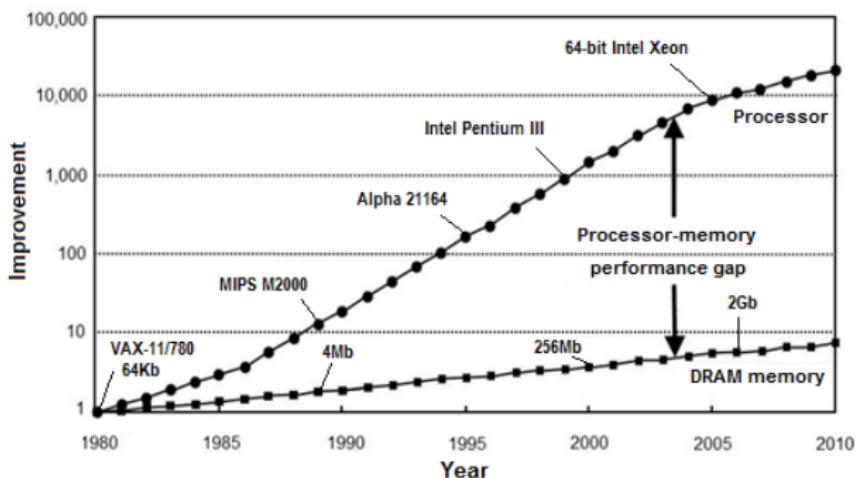
- ▶ CPUs use *registers* to store integers
  - ▶ 32-bit CPU: (most) registers are 32 bits
  - ▶ 64-bit CPU: (most) registers are 64 bits
- ▶ Address register size determines amount of addressable RAM
  - ▶ 20 bits (old 8086):  $2^{20}$  bytes addressable = 1 Mb
  - ▶ 32 bits:  $2^{32}$  bytes = 4 Gibi
    - ▶ **4 Gb limit** on 32-bit machines
  - ▶ 48 bits:  $2^{48}$  bytes = 256 Tebi
  - ▶ 64 bits:  $2^{64}$  bytes = 16 Exabytes (16 billion Gb)
- ▶ Most systems have switched to 64-bit
  - ▶ Not straightforward — need 64-bit OS
  - ▶ Similar to change from 16 to 32 bits in early 1990's

# Speed

- ▶ Memory speed: more important than most people realize
- ▶ CPU demands items from memory — often
  - ▶ ML instructions
  - ▶ Data
- ▶ 1Ghz CPU clock: 1 nanosecond per cycle
  - ▶ That's a fairly **slow** CPU
- ▶ Typical memory latency: 10 nanoseconds
  - ▶ If you have **fast** memory
- ▶ CPU waits **tens of cycles** for memory
  - ▶ Problem: your CPU slowed to 10% of its top speed

# Speed Gap

Figure from "A Survey of Different Approaches for Overcoming the Processor-Memory Bottleneck",  
by Efusheva et al., IJCSIT Vol 9, No 2, April 2017:



CPU speeds are improving faster than memory speeds

- ▶ CPU: around 20% improvement per year since 2004
- ▶ Memory: around 8% improvement per year

# Memory hierarchy

- ▶ We can build extremely fast memory
  - ▶ ... but the cost per byte is very high
- ▶ We can build extremely cheap memory
  - ▶ ... but it is very slow

# Memory hierarchy

- ▶ We can build extremely fast memory
  - ▶ ... but the cost per byte is very high
- ▶ We can build extremely cheap memory
  - ▶ ... but it is very slow
- ▶ Memory hierarchy: mix a variety of speeds and capacities
  - ▶ Faster memory serves as *cache* for slower memory

# Memory hierarchy

- ▶ We can build extremely fast memory
  - ▶ ... but the cost per byte is very high
- ▶ We can build extremely cheap memory
  - ▶ ... but it is very slow
- ▶ Memory hierarchy: mix a variety of speeds and capacities
  - ▶ Faster memory serves as *cache* for slower memory

## Modern memory hierarchy

Cost	Type	Speed
\$\$\$\$\$	CPU registers	fast!!!
\$\$\$	Memory on CPU (L1 Cache)	fast!!
\$\$	Memory on CPU (L2 Cache)	fast!
\$	L3 Cache	fast
¢	RAM	slow
	Disk	slow!!!!

# Standardized buses

- ▶ Allow devices to be “portable” to different motherboards
- ▶ Not all devices are integrated
- ▶ Common for extras, like
  - ▶ Modems
  - ▶ Network Interfaces
  - ▶ Sound
  - ▶ Game controllers
  - ▶ Custom devices
- ▶ Also common for high-performance items
  - ▶ Graphics
  - ▶ Drives

## General-purpose bus types: internal

- ▶ ISA bus (Industry Standard Architecture) 1984
  - ▶ Obsolete with Pentium
  - ▶ Included in some later motherboards
- ▶ PCI bus (Peripheral Component Interconnect) 1993
  - ▶ Introduced around Pentium
  - ▶ Better support for “plug-and-play”
- ▶ PCMCIA
  - ▶ Like PCI, but for laptops
  - ▶ Hot swappable
    - ▶ Can add devices while running
    - ▶ Can remove devices while running
- ▶ PCI-X (PCI eXtended) 1998
- ▶ PCIe (PCI express) 2004

## General-purpose bus types: external

- ▶ USB 1.0 (Universal Serial Bus) 1996
  - ▶ Serial: one bit transferred at a time
  - ▶ 12 Mbit/s max transmission rate
- ▶ USB 2.0 2000
  - ▶ 480 Mbit/s max transmission rate
- ▶ USB 3.0 2008
  - ▶ 5 Gbit/s max transmission rate
  - ▶ USB 3.1, 3.2 have faster transmission rates
- ▶ USB 4 2019
  - ▶ Requires USB-C connectors
  - ▶ Carries other protocols (USB 3.2, DisplayPort, PCIe)
- ▶ FireWire 1995 – 2008
- ▶ Thunderbolt 2011

# Graphics bus types

- ▶ VESA Local Bus 1992 — 1993
  - ▶ Video Electronics Standards Association
  - ▶ Closely tied to Intel 80486 CPU memory bus design
  - ▶ ISA bus was becoming bottleneck for graphics
  - ▶ Short-term solution: add an extra slot to an ISA slot
    - ▶ Could use ordinary ISA card and ignore extra
  - ▶ Obsolete with PCI
- ▶ AGP (Accelerated Graphics Port) 1997 — 2004
  - ▶ Introduced when PCI became bottleneck for graphics
  - ▶ Obsolete with PCIe

# Buses for drives

- ▶ IDE (Integrated Drive Electronics), same as (parallel) ATA: IBM AT computer Attachment
  - ▶ Each IDE interface can control at most 2 drives (with 1 cable)
    - ▶ Device 0: “primary” or “master”; normally “C” drive
    - ▶ Device 1: “secondary” or “slave”; normally “D” drive
    - ▶ Historical “master and slave” terminology has seen controversy
  - ▶ Controllers typically integrated into motherboard
  - ▶ Uses “ribbon cables” (40 or 80 wires)



An ancient 40-wire ribbon cable, for floppy drives

## Buses for drives, part 2

- ▶ SATA (Serial ATA) 2003
  - ▶ Gradually replaced IDE
  - ▶ Cables have 7 wires
  - ▶ Each cable connects 1 device
- ▶ SCSI (Small Computer System Interface) 1982
  - ▶ Original interface was parallel (8, then 16 bits)
- ▶ SAS (Serial Attached SCSI) 2004

Units  
oooooo

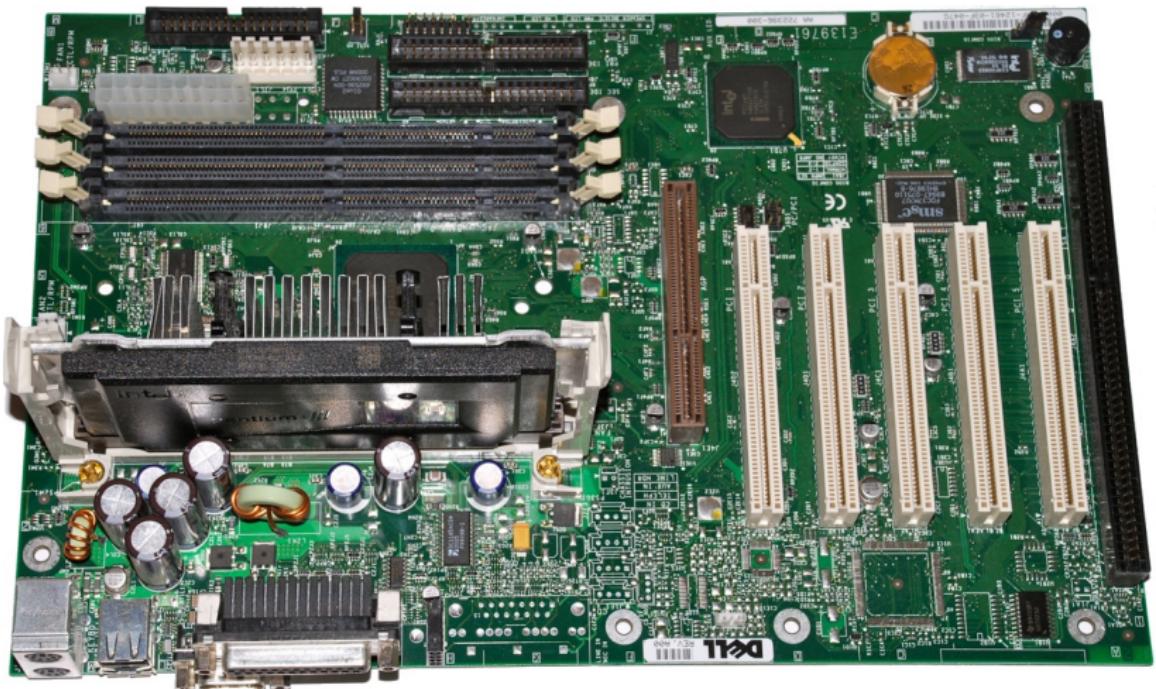
Organization  
oooooooo

CPU  
ooooooo

RAM  
ooooo

Buses  
ooooooo●

VMs  
oooooooooooo



Units  
oooooo

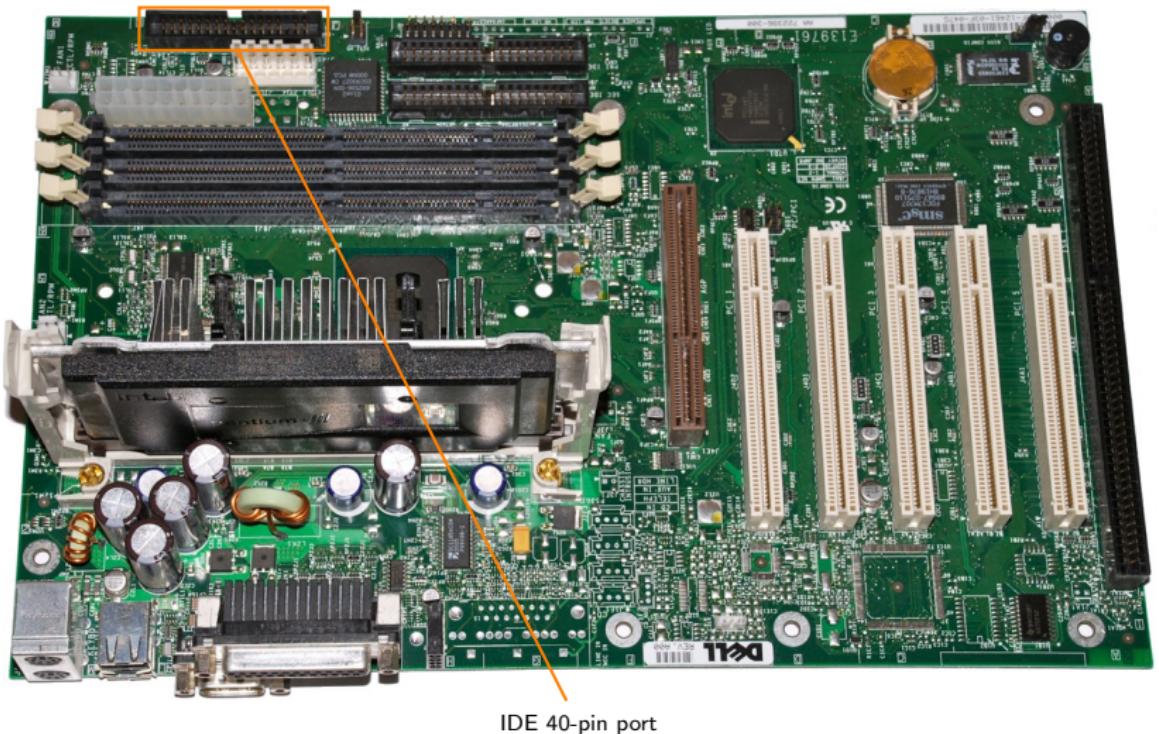
Organization  
oooooooo

CPU  
ooooooo

RAM  
ooooo

Buses  
ooooooo●

VMs  
oooooooooooo



IDE 40-pin port

Units  
oooooo

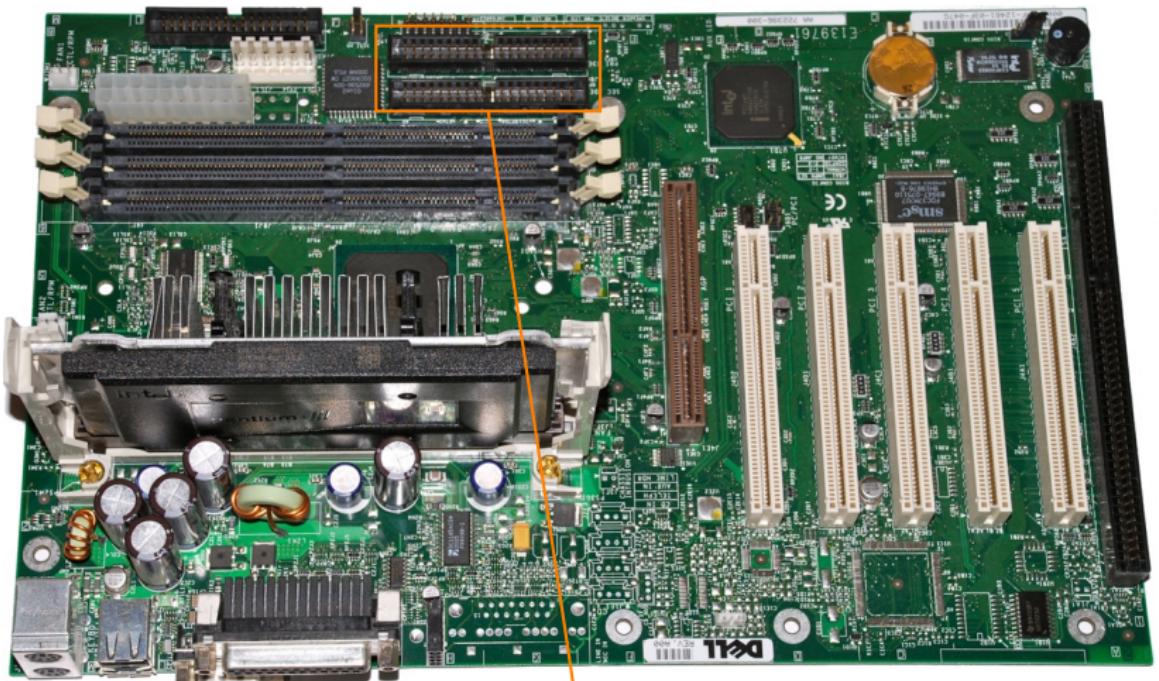
Organization  
oooooooo

CPU  
ooooooo

RAM  
ooooo

Buses  
ooooooo●

VMs  
oooooooooooo



IDE 80-pin ports

Units  
oooooo

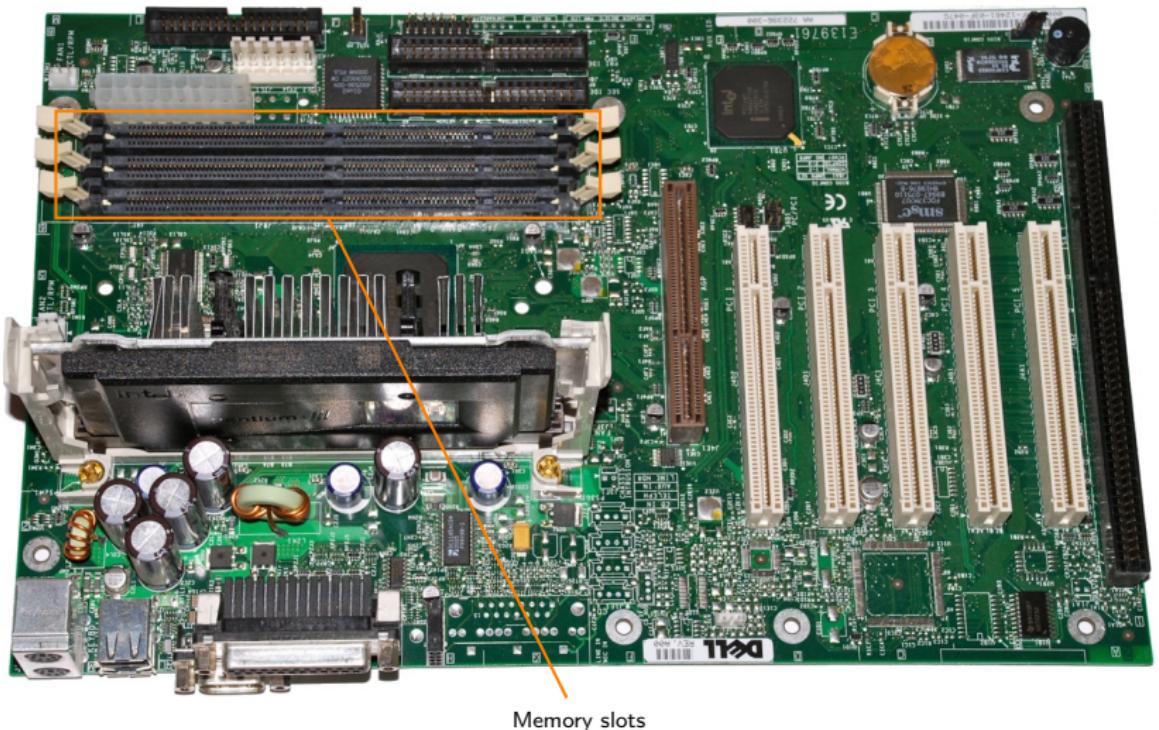
Organization  
oooooooo

CPU  
ooooooo

RAM  
ooooo

Buses  
ooooooo●

VMs  
oooooooooooo



Units  
oooooo

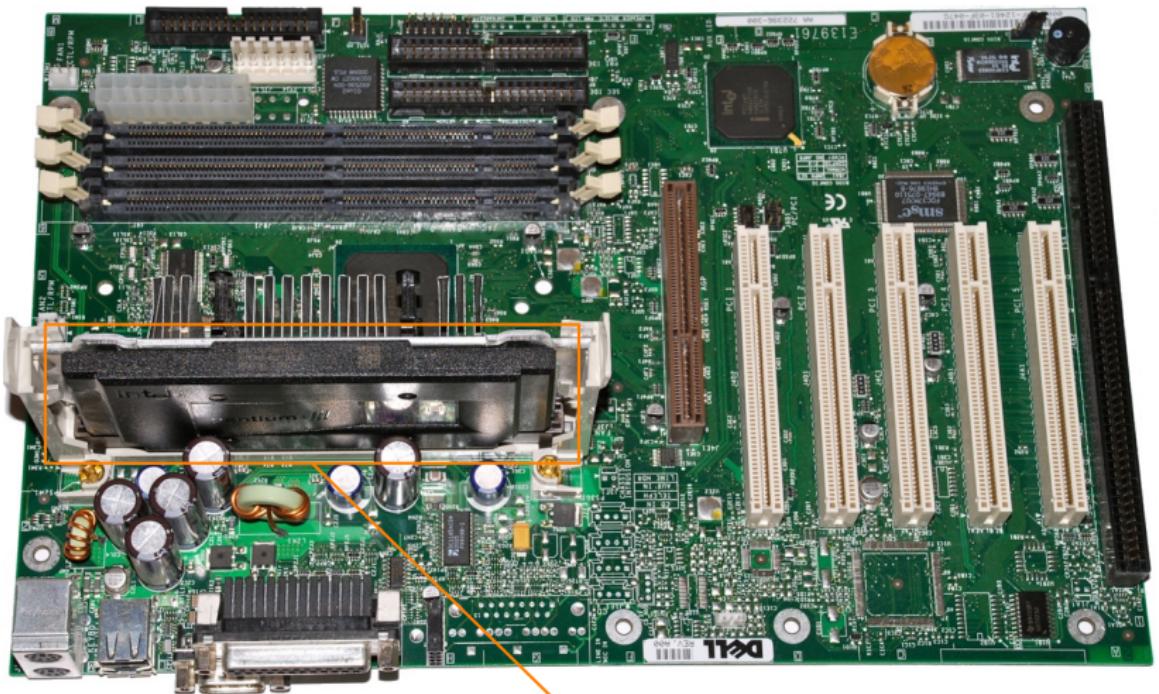
Organization  
oooooooo

CPU  
ooooooo

RAM  
ooooo

Buses  
ooooooo●

VMs  
oooooooooooo



Pentium 3 CPU (vertical)

Units  
oooooo

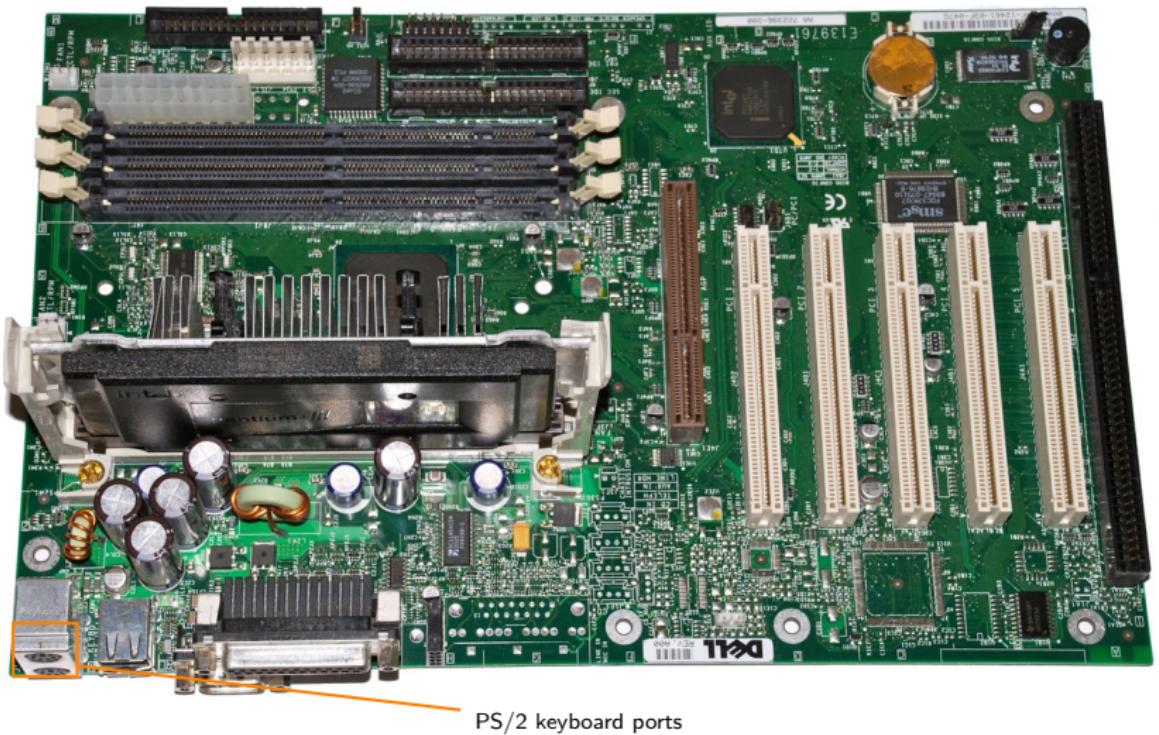
Organization  
oooooooo

CPU  
ooooooo

RAM  
ooooo

Buses  
ooooooo●

VMs  
oooooooooooo



PS/2 keyboard ports

Units  
oooooo

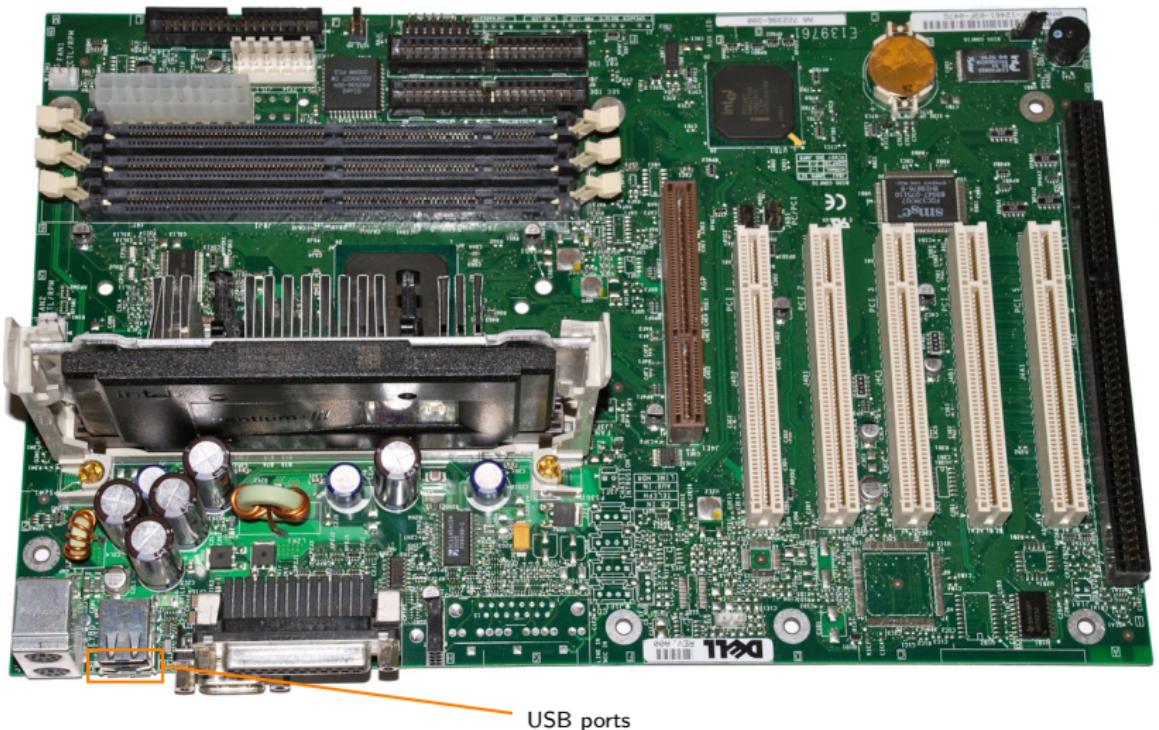
Organization  
oooooooo

CPU  
ooooooo

RAM  
ooooo

Buses  
ooooooo●

VMs  
oooooooooooo



Units  
oooooo

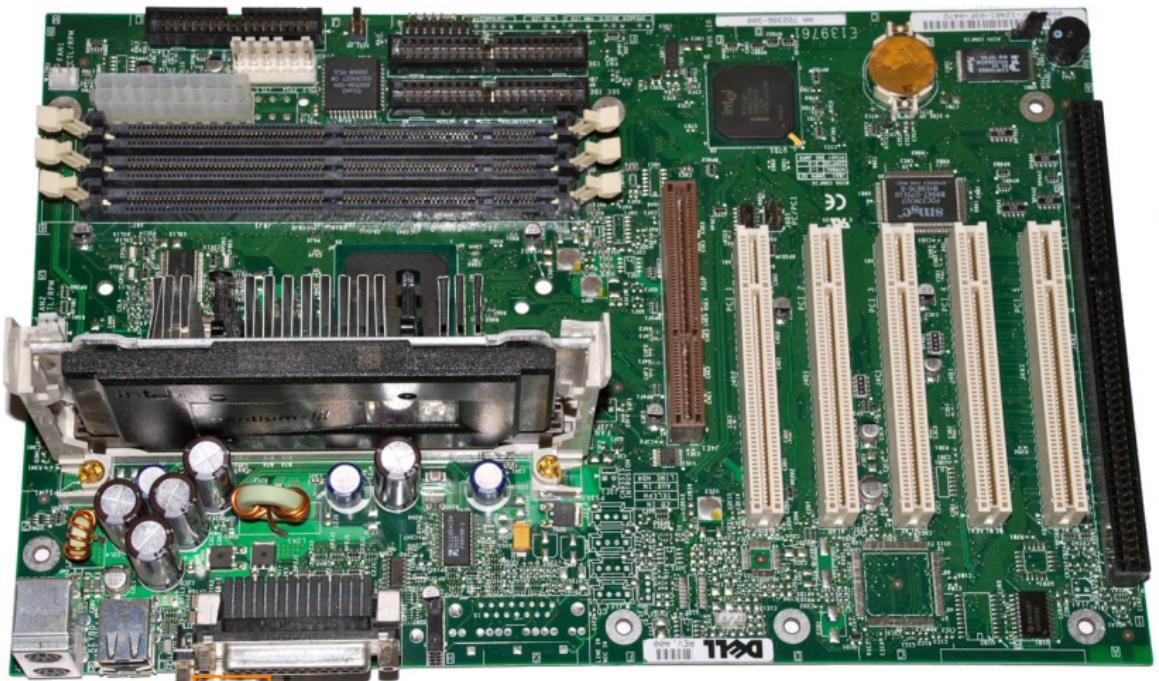
Organization  
oooooooo

CPU  
ooooooo

RAM  
ooooo

Buses  
ooooooo●

VMs  
oooooooooooo



Serial port

Units  
oooooo

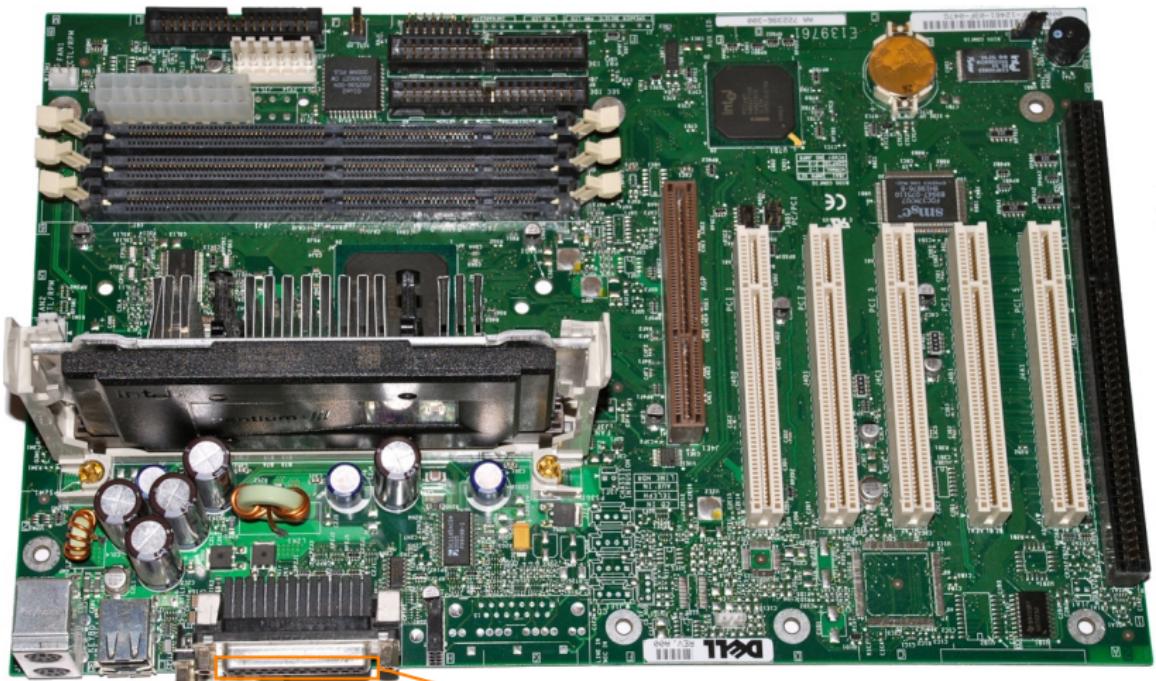
Organization  
oooooooo

CPU  
ooooooo

RAM  
ooooo

Buses  
ooooooo●

VMs  
oooooooooooo



Parallel port

Units  
oooooo

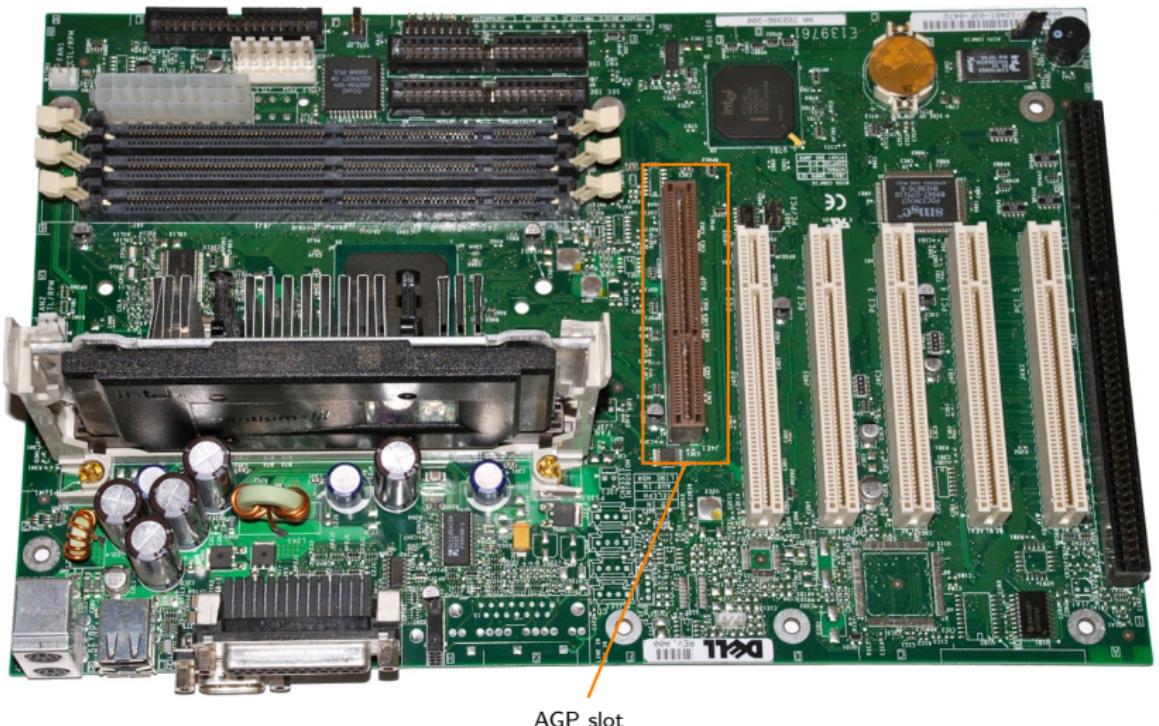
Organization  
oooooooo

CPU  
ooooooo

RAM  
ooooo

Buses  
ooooooo●

VMs  
oooooooooooo



Units  
oooooo

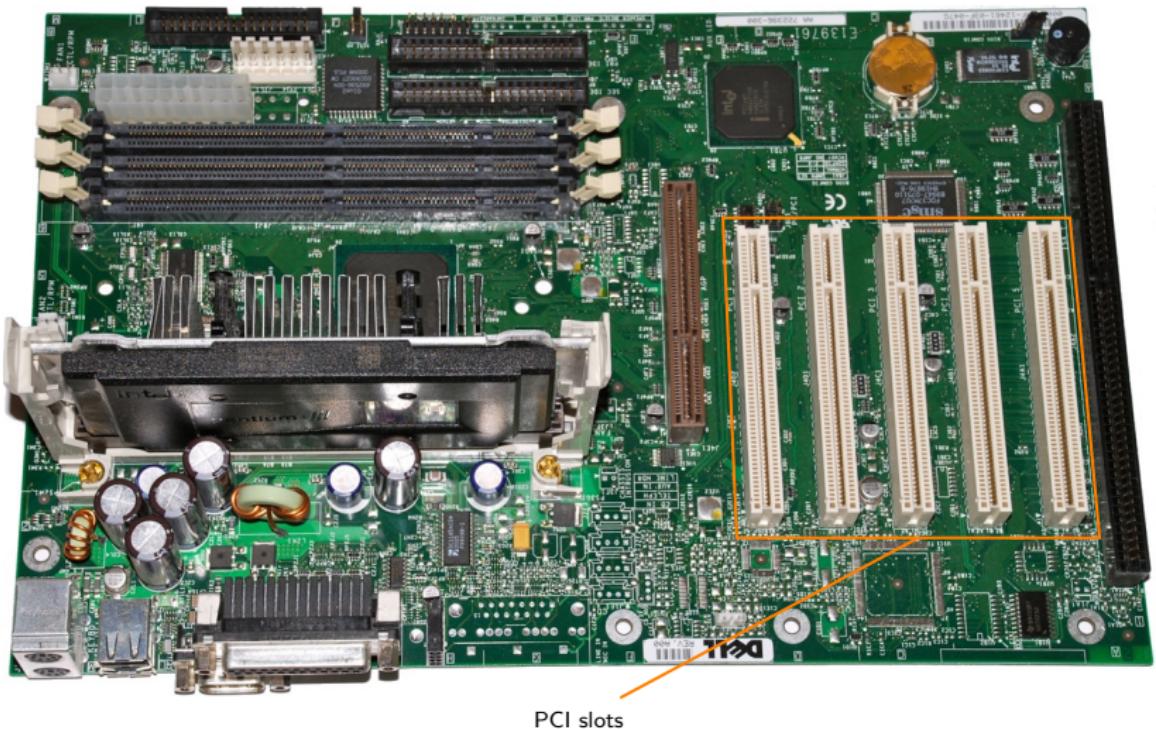
Organization  
oooooooo

CPU  
ooooooo

RAM  
ooooo

Buses  
ooooooo●

VMs  
oooooooooooo



Units  
oooooo

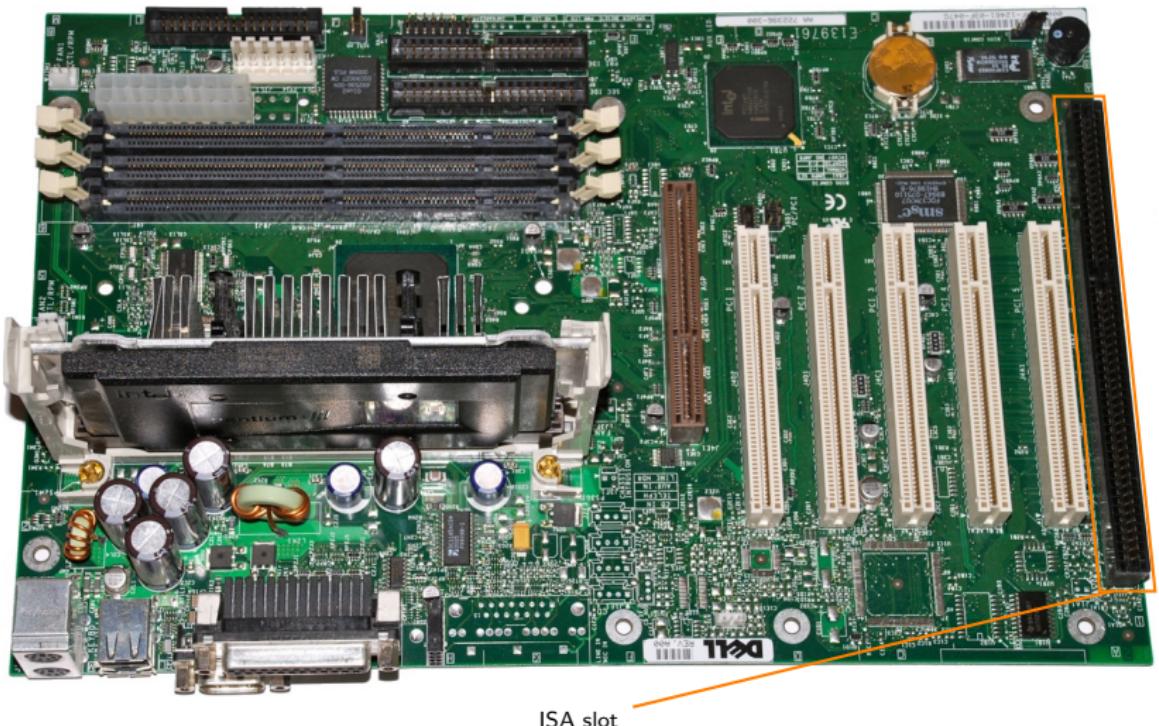
Organization  
oooooooo

CPU  
ooooooo

RAM  
ooooo

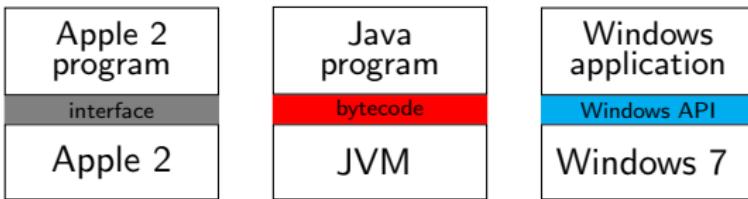
Buses  
ooooooo●

VMs  
oooooooooooo



# Virtual machines

- ▶ All software has an *interface* to an *environment*
- ▶ “Environment” can be hardware, software, or both:



- ▶ The software can run *anywhere* the interface is provided
- ▶ *Virtual machine*: software that provides a machine interface
  - ▶ **Host** refers to the “actual system” beneath the VM
  - ▶ **Guest** refers to anything running on the VM
  - ▶ More than one guest can run on a single host
  - ▶ *Virtualization*: same hardware; utilizes CPU support
  - ▶ *Emulation*: different hardware; uses software to execute

# Virtual devices

Virtual machine software must provide *virtual devices*

- ▶ Virtual display
  - ▶ VM display can appear in a window
- ▶ Virtual keyboard, mouse
  - ▶ Allows keyboard, mouse to be used in VM
  - ▶ Allows keyboard, mouse to be used *outside* VM
- ▶ Virtual drives
  - ▶ Guest drives are (large) files on host system
  - ▶ CD/DVD drives can connect to disk images or actual drives
- ▶ Virtual network
  - ▶ Allows VMs to connect to host machine's network
  - ▶ Allows VMs to connect to host machine
  - ▶ Allows VMs to connect to each other
- ▶ Other devices
  - ▶ Can choose to connect, disconnect USB devices to VM

Units  
oooooo

Organization  
oooooooo

CPU  
oooooooo

RAM  
oooooo

Buses  
oooooooo

VMs  
oo●oooooooo

# VirtualBox: x86 virtualization software by Oracle

- ▶ <http://www.virtualbox.org>
- ▶ Homeworks will be designed for this
  - ▶ May be possible to import into VMWare — have not tried this
- ▶ Open Source: Released under GPL version 2
- ▶ Runs on Windows, Linux, Macintosh, Solaris **host OS's**
- ▶ Has support for a variety of **guest OS's**

Units  
oooooo

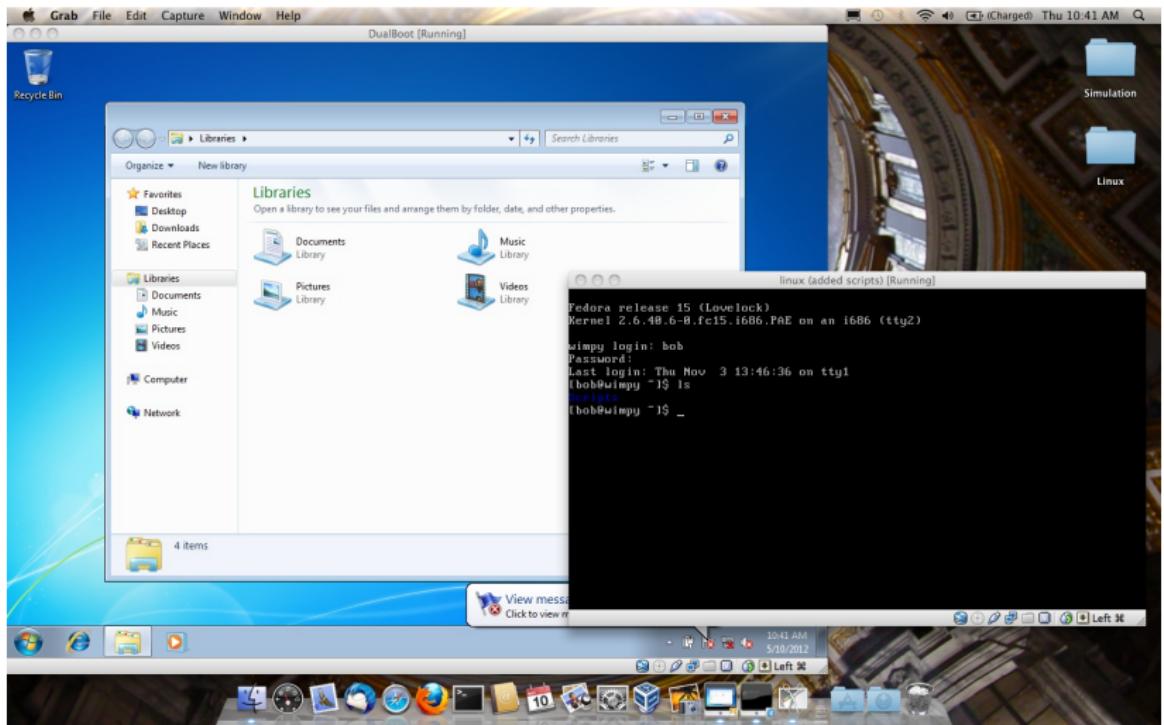
Organization  
oooooooo

CPU  
oooooooo

RAM  
oooooo

Buses  
oooooooo

VMs  
oooo●oooooooo



Screen capture showing Windows and Linux VMs running in VirtualBox Mac OS

Units  
oooooo

Organization  
oooooooo

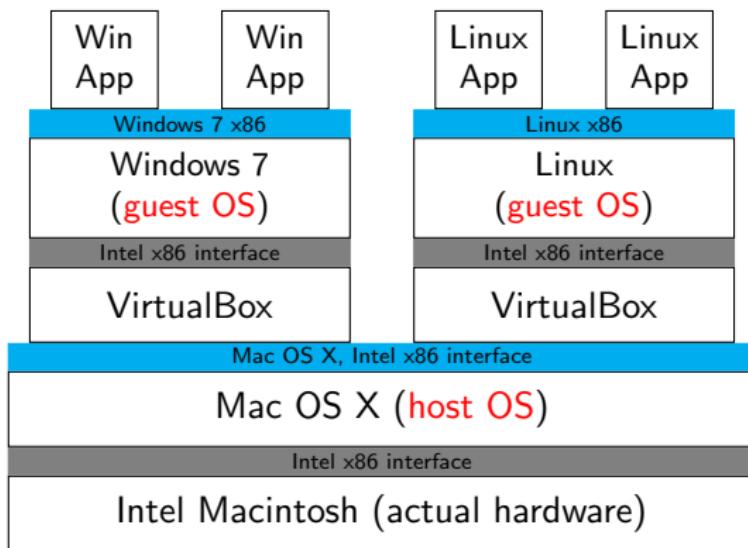
CPU  
oooooooo

RAM  
oooooo

Buses  
oooooooo

VMs  
oooo●oooooooo

## Discussion of screenshot



Units  
oooooo

Organization  
oooooooo

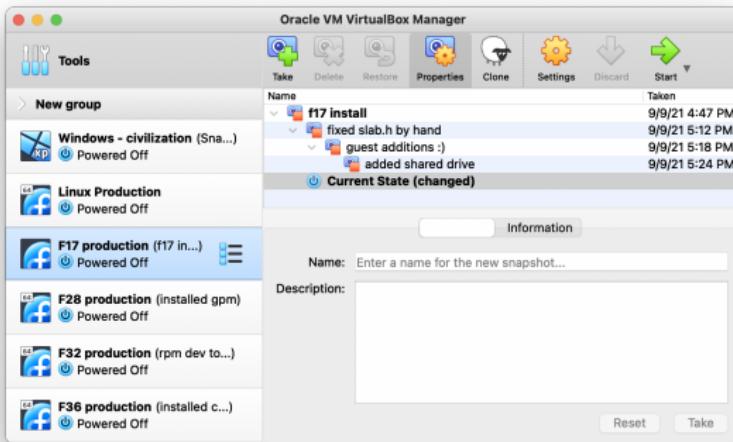
CPU  
oooooooo

RAM  
oooooo

Buses  
oooooooo

VMs  
oooooooooooo

# Using VirtualBox



Units  
oooooo

Organization  
oooooooo

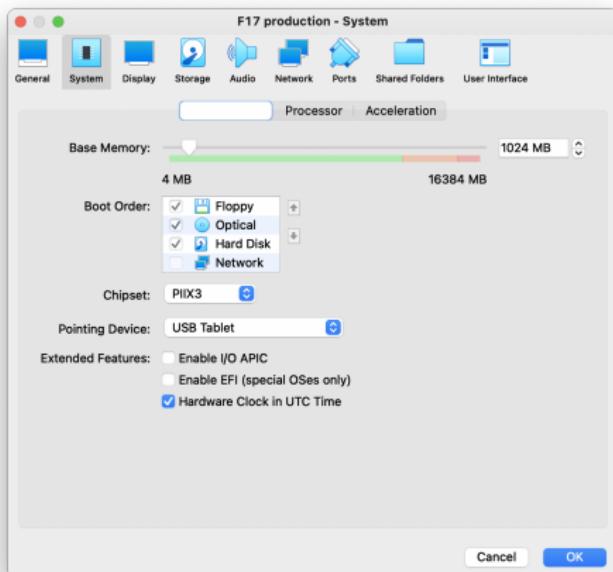
CPU  
oooooooo

RAM  
oooo

Buses  
oooooooo

VMs  
oooooooo●ooo

# VM Settings: System



Units  
oooooo

Organization  
oooooooo

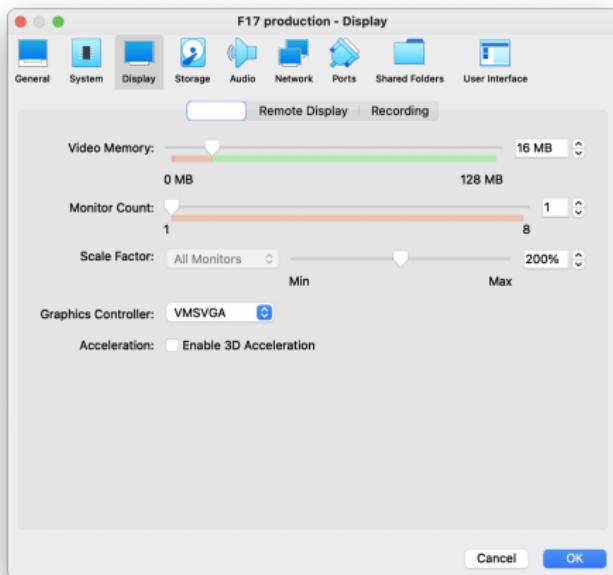
CPU  
oooooooo

RAM  
oooooo

Buses  
oooooooo

VMs  
oooooooo●○○

# VM Settings: Display



Units  
oooooo

Organization  
oooooooo

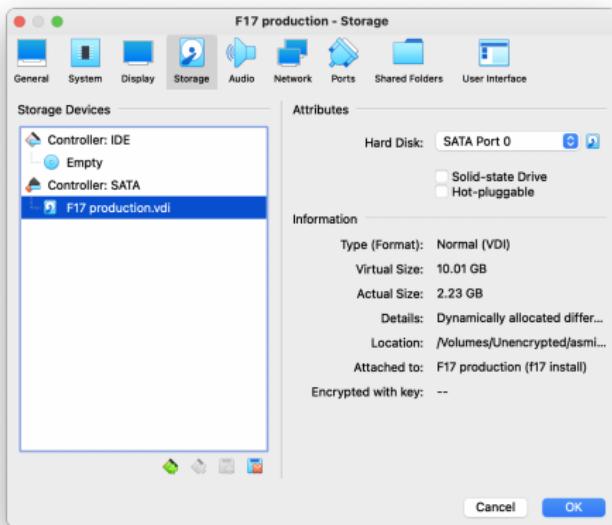
CPU  
oooooooo

RAM  
oooooo

Buses  
oooooooo

VMs  
oooooooooooo●

# VM Settings: Storage



Units  
oooooo

Organization  
oooooooo

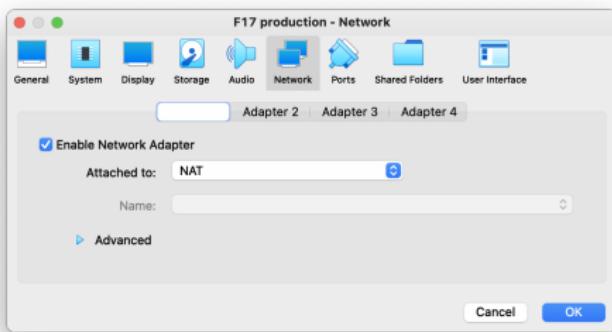
CPU  
oooooooo

RAM  
oooooo

Buses  
oooooooo

VMs  
oooooooooooo●

# VM Settings: Network



## An appropriate xkcd comic: <http://xkcd.com/394>

THERE'S BEEN A LOT OF CONFUSION OVER 1024 VS 1000,  
KBYTE VS KBIT, AND THE CAPITALIZATION FOR EACH.

HERE, AT LAST, IS A SINGLE, DEFINITIVE STANDARD:

SYMBOL	NAME	SIZE	NOTES
kB	KILOBYTE	1024 BYTES OR 1000 BYTES	1000 BYTES DURING LEAP YEARS, 1024 OTHERWISE
KB	KELLY-BOOTLE STANDARD UNIT	1012 BYTES	COMPROMISE BETWEEN 1000 AND 1024 BYTES
KiB	IMAGINARY KILOBYTE	1024 $\sqrt{F}$ BYTES	USED IN QUANTUM COMPUTING
kb	INTEL KILOBYTE	1023.937528 BYTES	CALCULATED ON PENTIUM FPU.
Kb	DRIVEMAKER'S KILOBYTE	CURRENTLY 908 BYTES	SHRINKS BY 4 BYTES EACH YEAR FOR MARKETING REASONS
KBa	BAKER'S KILOBYTE	1152 BYTES	9 BITS TO THE BYTE SINCE YOU'RE SUCH A GOOD CUSTOMER

Units  
oooooo

Organization  
oooooooo

CPU  
oooooooo

RAM  
oooooo

Buses  
oooooooo

VMs  
oooooooooooo

# End of lecture