

# Internet Protocol

ComS 252 — Iowa State University

Barry Britt and Andrew Miner

# What is IPv4?

- ▶ Version 4 of the Internet Protocol
- ▶ This is at the “Network Layer” in the OSI model
  - ▶ Deals with sending packets from one machine to another
  - ▶ Packets may go through several **routers**
- ▶ Let's see how it works

# IP Packets

IP Packets consist of

1. The packet **header**

- ▶ Like the envelope used to send postal mail
- ▶ Contains information to send the data
- ▶ Basically, a “struct” with several fields
- ▶ Protocol specifies:
  - ▶ Order of the fields in the header
  - ▶ Number of bits for each field

2. Actual data

- ▶ Interpreted based on the protocol used on top of IP
  - ▶ I.e., the protocol at the transport layer in the OSI model

# What is in a packet header?

- ▶ Version number (should be 4)
  - ▶ First 4 bits of the packet

# What is in a packet header?

- ▶ Version number (should be 4)
  - ▶ First 4 bits of the packet
- ▶ Total length (16 bits)
  - ▶ Total packet length, in bytes
  - ▶ Minimum is 20 bytes (20-byte header plus 0 bytes of data)
  - ▶ Maximum is

# What is in a packet header?

- ▶ Version number (should be 4)
  - ▶ First 4 bits of the packet
- ▶ Total length (16 bits)
  - ▶ Total packet length, in bytes
  - ▶ Minimum is 20 bytes (20-byte header plus 0 bytes of data)
  - ▶ Maximum is  $2^{16} - 1 = 65535$  bytes

# What is in a packet header?

- ▶ Version number (should be 4)
  - ▶ First 4 bits of the packet
- ▶ Total length (16 bits)
  - ▶ Total packet length, in bytes
  - ▶ Minimum is 20 bytes (20-byte header plus 0 bytes of data)
  - ▶ Maximum is  $2^{16} - 1 = 65535$  bytes
- ▶ Information for **fragmenting** packets
  - ▶ A packet may be broken into pieces and reassembled
  - ▶ Done so that **packets** will fit into **frames**

# What is in a packet header?

- ▶ Version number (should be 4)
  - ▶ First 4 bits of the packet
- ▶ Total length (16 bits)
  - ▶ Total packet length, in bytes
  - ▶ Minimum is 20 bytes (20-byte header plus 0 bytes of data)
  - ▶ Maximum is  $2^{16} - 1 = 65535$  bytes
- ▶ Information for **fragmenting** packets
  - ▶ A packet may be broken into pieces and reassembled
  - ▶ Done so that **packets** will fit into **frames**
- ▶ **TTL**: Time to Live (8 bits)
  - ▶ Specifies maximum lifetime of a packet
  - ▶ Keeps packets from going in circles indefinitely
  - ▶ Original spec was lifetime in **seconds**
  - ▶ In practice it is the maximum “hop count”:
    - ▶ Each router decrements the TTL by one
    - ▶ Routers discard packets with TTL field of 0



## What is in a packet header? (2)

- ▶ Protocol (8 bits)
  - ▶ What protocol is on top of this?
  - ▶ Can be **ICMP**, **TCP**, **UDP**, and others
  - ▶ ICMP: Internet Control Message Protocol
    - ▶ Used for query and error messages within IP

## What is in a packet header? (2)

- ▶ Protocol (8 bits)
  - ▶ What protocol is on top of this?
  - ▶ Can be **ICMP**, **TCP**, **UDP**, and others
  - ▶ ICMP: Internet Control Message Protocol
    - ▶ Used for query and error messages within IP
- ▶ Source address (32 bits)
- ▶ Destination address (32 bits)

## What is in a packet header? (2)

- ▶ Protocol (8 bits)
  - ▶ What protocol is on top of this?
  - ▶ Can be **ICMP**, **TCP**, **UDP**, and others
  - ▶ ICMP: Internet Control Message Protocol
    - ▶ Used for query and error messages within IP
- ▶ Source address (32 bits)
- ▶ Destination address (32 bits)
- ▶ Header Checksum
  - ▶ For detecting errors
- ▶ Other fields. . .
- ▶ IHL (4 bits): number of 32-bit words in the packet header
- ▶ Options: zero or more **optional** words (32 bits each)
  - ▶ Typically unused

# Addresses in IPv4

- ▶ 32 bits long
  - ▶ Number of **unique** addresses in IPv4:

# Addresses in IPv4

- ▶ 32 bits long
  - ▶ Number of **unique** addresses in IPv4:  $2^{32} \approx 4.29$  billion
  - ▶ Number of devices “on the Internet” using IPv4:

# Addresses in IPv4

- ▶ 32 bits long
  - ▶ Number of **unique** addresses in IPv4:  $2^{32} \approx 4.29$  billion
  - ▶ Number of devices “on the Internet” using IPv4:  $> 2^{32}$ 
    - ▶ Total number of devices on the Internet:  $\approx 5$  billion in 2010
  - ▶ How is this possible?

# Addresses in IPv4

- ▶ 32 bits long
  - ▶ Number of **unique** addresses in IPv4:  $2^{32} \approx 4.29$  billion
  - ▶ Number of devices “on the Internet” using IPv4:  $> 2^{32}$ 
    - ▶ Total number of devices on the Internet:  $\approx 5$  billion in 2010
  - ▶ How is this possible?
    - ▶ There must be lots of devices using the **same** address
    - ▶ We will see why and how this works, later

# Addresses in IPv4

- ▶ 32 bits long
  - ▶ Number of **unique** addresses in IPv4:  $2^{32} \approx 4.29$  billion
  - ▶ Number of devices “on the Internet” using IPv4:  $> 2^{32}$ 
    - ▶ Total number of devices on the Internet:  $\approx 5$  billion in 2010
  - ▶ How is this possible?
    - ▶ There must be lots of devices using the **same** address
    - ▶ We will see why and how this works, later
- ▶ Written as 4 integers separated by dots
  - ▶ Each integer is 8 bits
  - ▶ For example:
$$192.168.2.17 = 11000000.10101000.00000010.00010001$$
- ▶ Governed by **IANA** (Internet Assigned Numbers Authority)
  - ▶ A department of **ICANN**  
(Internet Corporation for Assigned Names and Numbers)



# Subnetworks

How are subnetworks handled in IPv4?

# Subnetworks

How are subnetworks handled in IPv4?

- ▶ 1981 — 1993: **Classful network**
  - ▶ Introduced by [RFC 791](#) — Internet Protocol
  - ▶ Splits the address space into **classes**
    - ▶ Class A, Class B, Class C, Class D, Class E
    - ▶ Classes D and E were not fully defined
- ▶ 1993 — today: **Classless Inter-Domain Routing**
  - ▶ Introduced by [RFC 1518](#) and [RFC 1519](#)

# Class A: Huge subnetworks

- ▶ First bit of 32-bit address is 0
- ▶ Next 7 bits: the subnetwork ID
  - ▶ At most  $2^7 = 128$  Class A networks
- ▶ Final 24 bits: host ID within the subnetwork
  - ▶ At most  $2^{24} = 16,777,216$  hosts in the subnet
- ▶ IP address has the form  
0nnnnnnn.HHHHHHHH.HHHHHHHH.HHHHHHHH

Example:

- ▶ Address 10.9.8.7 =  
00001010.00001001.00001000.00000111
- ▶ Class A subnetwork: 10.
- ▶ Host within subnetwork: 9.8.7

## Class B: Large subnetworks

- ▶ First bit of 32-bit address is **1**
  - ▶ Because it is not Class A
- ▶ Second bit of 32-bit address is **0**
- ▶ Next 14 bits: the subnetwork ID
  - ▶ At most  $2^{14} = 16,384$  Class B networks
- ▶ Final 16 bits: host ID within the subnetwork
  - ▶ At most  $2^{16} = 65,536$  hosts in the subnet
- ▶ IP address has the form  
**10**nnnnnn . nnnnnnnn . HHHHHHHH . HHHHHHHH

# Class C

- ▶ First bit of 32-bit address is **1**
  - ▶ Because it is not Class A
- ▶ Second bit of 32-bit address is **1**
  - ▶ Because it is not Class B
- ▶ Third bit of 32-bit address is **0**
- ▶ Next 21 bits: the subnetwork ID
  - ▶ At most  $2^{21} = 2,097,152$  Class C networks
- ▶ Final 8 bits: host ID within the subnetwork
  - ▶ At most  $2^8 = 256$  hosts in the subnet
- ▶ IP address has the form  
**110****nnnnn**.**nnnnnnnn**.**nnnnnnnn**.HHHHHHHH

# Classes D and E

## Class D

- ▶ First 4 bits are **1110**
- ▶ IP address has the form  
**1110**xxxx . xxxxxxxx . xxxxxxxx . xxxxxxxx

## Class E

- ▶ First 4 bits are **1111**
- ▶ IP address has the form  
**1111**xxxx . xxxxxxxx . xxxxxxxx . xxxxxxxx

# Why the switch away from Classful network in 1993?

# Why the switch away from Classful network in 1993?

- ▶ Too “wasteful” of IP addresses
- ▶ E.g., suppose I want a subnet with 500 hosts
  - ▶ I need a **Class B** subnet with 65,536 addresses
- ▶ E.g., suppose I want a subnet with 70,000 hosts
  - ▶ I need a **Class A** subnet with 16 million addresses



# Why the switch away from Classful network in 1993?

- ▶ Too “wasteful” of IP addresses
- ▶ E.g., suppose I want a subnet with 500 hosts
  - ▶ I need a **Class B** subnet with 65,536 addresses
- ▶ E.g., suppose I want a subnet with 70,000 hosts
  - ▶ I need a **Class A** subnet with 16 million addresses
- ▶ Classless Inter–Domain Routing is more flexible. . .

# Classless Inter-Domain Routing (CIDR)

- ▶ Every IP address is **partitioned** as follows:
  - ▶ The first  $n$  bits are the subnet
  - ▶ The last  $32 - n$  bits are the host name within the subnet
  - ▶  $n$  may vary by subnet
- ▶ A subnet name may be specified by giving
  1. The number of prefix bits  $n$  for the subnet
  2. A complete IP address with last  $32 - n$  bits equal to zero
- ▶  $n$  can be specified after the address, e.g.:  
192.168.2.0/24
- ▶  $n$  can be specified by giving the **subnet mask**, e.g.:  
192.168.2.0/255.255.255.0
- ▶ The suffix of **all zeroes** is the subnet name
- ▶ The suffix of **all ones** is the **broadcast** address

## 16-bit subnet prefix example

Subnet mask: 255.255.0.0 = 11111111.11111111.00000000.00000000

Network: 129.186.0.0 = 10000001.10111010.00000000.00000000

First usable: 129.186.0.1 = 10000001.10111010.00000000.00000001

Last usable: 129.186.255.254 = 10000001.10111010.11111111.11111110

Broadcast: 129.186.255.255 = 10000001.10111010.11111111.11111111

## 16-bit subnet prefix example

Subnet mask: 255.255.0.0 = 11111111.11111111.00000000.00000000  
Network: 129.186.0.0 = 10000001.10111010.00000000.00000000  
First usable: 129.186.0.1 = 10000001.10111010.00000000.00000001  
Last usable: 129.186.255.254 = 10000001.10111010.11111111.11111110  
Broadcast: 129.186.255.255 = 10000001.10111010.11111111.11111111

## 24-bit subnet prefix example

Subnet mask: 255.255.255.0 = 11111111.11111111.11111111.00000000  
Network: 192.168.2.0 = 11000000.10101000.00000010.00000000  
First usable: 192.168.2.1 = 11000000.10101000.00000010.00000001  
Last usable: 192.168.2.254 = 11000000.10101000.00000010.11111110  
Broadcast: 192.168.2.255 = 11000000.10101000.00000010.11111111

## 16-bit subnet prefix example

Subnet mask: 255.255.0.0 = 11111111.11111111.00000000.00000000  
Network: 129.186.0.0 = 10000001.10111010.00000000.00000000  
First usable: 129.186.0.1 = 10000001.10111010.00000000.00000001  
Last usable: 129.186.255.254 = 10000001.10111010.11111111.11111110  
Broadcast: 129.186.255.255 = 10000001.10111010.11111111.11111111

## 24-bit subnet prefix example

Subnet mask: 255.255.255.0 = 11111111.11111111.11111111.00000000  
Network: 192.168.2.0 = 11000000.10101000.00000010.00000000  
First usable: 192.168.2.1 = 11000000.10101000.00000010.00000001  
Last usable: 192.168.2.254 = 11000000.10101000.00000010.11111110  
Broadcast: 192.168.2.255 = 11000000.10101000.00000010.11111111

## 28-bit subnet prefix example

Subnet mask: 255.255.255.240 = 11111111.11111111.11111111.11110000  
Network: 172.1.1.176 = 10101100.00000001.00000001.10110000  
First usable: 172.1.1.177 = 10101100.00000001.00000001.10110001  
Last usable: 172.1.1.190 = 10101100.00000001.00000001.10111110  
Broadcast: 172.1.1.191 = 10101100.00000001.00000001.10111111

# Relationship between Classful and Classless

## Old Class A :

- ▶ Subnets 0.0.0.0/8 — 127.0.0.0/8
- ▶ Subnet mask is 255.0.0.0

## Old Class B :

- ▶ Subnets 128.0.0.0/16 — 191.255.0.0/16
- ▶ Subnet mask is 255.255.0.0

## Old Class C :

- ▶ Subnets 192.0.0.0/24 — 223.255.255.0/24
- ▶ Subnet mask is 255.255.255.0

# Private Networks

- ▶ Suppose I want to set up an IP network at home
- ▶ Network sits behind a gateway plus router plus firewall
- ▶ No reason to connect to any of my machines from “outside”
- ▶ Do I need to contact IANA for IP addresses?

# Private Networks

- ▶ Suppose I want to set up an IP network at home
- ▶ Network sits behind a gateway plus router plus firewall
- ▶ No reason to connect to any of my machines from “outside”
- ▶ Do I need to contact IANA for IP addresses? **NO**
- ▶ Can I use whatever IP addresses I want?



# Private Networks

- ▶ Suppose I want to set up an IP network at home
- ▶ Network sits behind a gateway plus router plus firewall
- ▶ No reason to connect to any of my machines from “outside”
- ▶ Do I need to contact IANA for IP addresses? **NO**
- ▶ Can I use whatever IP addresses I want? **YES** but ...

# Private Networks

- ▶ Suppose I want to set up an IP network at home
- ▶ Network sits behind a gateway plus router plus firewall
- ▶ No reason to connect to any of my machines from “outside”
- ▶ Do I need to contact IANA for IP addresses? **NO**
- ▶ Can I use whatever IP addresses I want? **YES** but ...
  - ▶ Suppose I choose an address that corresponds to google.com
  - ▶ When I try to send packets to google.com they will instead go to one of my machines


# Private Networks

- ▶ Suppose I want to set up an IP network at home
- ▶ Network sits behind a gateway plus router plus firewall
- ▶ No reason to connect to any of my machines from “outside”
- ▶ Do I need to contact IANA for IP addresses? **NO**
- ▶ Can I use whatever IP addresses I want? **YES** but ...
  - ▶ Suppose I choose an address that corresponds to `google.com`
  - ▶ When I try to send packets to `google.com` they will instead go to one of my machines
- ▶ So, what addresses can I use?

# Private Network Addresses

- ▶ [RFC 1918](#) sets aside blocks of addresses for **private use**
- ▶ One “Class A” subnet: 10.
  - ▶ Addresses 10.0.0.0 — 10.255.255.255
- ▶ 16 “Class B” subnets: 172.16 — 172.31
  - ▶ Addresses 172.16.0.0 — 172.31.255.255
- ▶ 256 “Class C” subnets: 192.168.0 — 192.168.255
  - ▶ Addresses 192.168.0.0 — 192.168.255.255
- ▶ These addresses are **guaranteed** to remain private
- ▶ You may use these addresses without contacting IANA
  - ▶ That's why they are set aside
- ▶ Using CIDR, you may use whatever subnets you want
  - ▶ I can use 10.8.6.4/30 for a private subnet if I want ...

# Example Home LAN




Suppose gypsy is a typical router

- ▶ Packets from one subnet are forwarded to the other

But what happens when I run a browser on **cambot**?

# Example Home LAN




Suppose gypsy is a typical router

- ▶ Packets from one subnet are forwarded to the other

But what happens when I run a browser on cambot?

1. cambot (192.168.42.100) sends packets to httpd running on remote server (say, 129.186.23.166)
2. How do packets **get back** to 192.168.42.100?
  - ▶ Private IP address — cannot route packets to it

# Example Home LAN



Suppose gypsy is a typical router

- ▶ Packets from one subnet are forwarded to the other

But what happens when I run a browser on cambot?

1. cambot (192.168.42.100) sends packets to  
httpd running on remote server (say, 129.186.23.166)
2. How do packets **get back** to 192.168.42.100?
  - ▶ Private IP address — cannot route packets to it

gypsy cannot be an “ordinary” router...

# NAT: Network Address Translation

A.k.a. “IP Masquerading”

- ▶ We have a gateway router where
  - ▶ One side is a private network
  - ▶ One side is a “real” address
- ▶ Packets coming from the private network:
  - ▶ Rewrite the packet header
  - ▶ Pretend the packet came from the gateway’s “real” address
  - ▶ Remember that we did this
- ▶ Packets coming from the other end:
  - ▶ If this is a response to a modified packet, then rewrite the header and send it to the right private IP address
  - ▶ Otherwise, no change



# NAT: Network Address Translation

A.k.a. “IP Masquerading”

- ▶ We have a gateway router where
  - ▶ One side is a private network
  - ▶ One side is a “real” address
- ▶ Packets coming from the private network:
  - ▶ Rewrite the packet header
  - ▶ Pretend the packet came from the gateway’s “real” address
  - ▶ Remember that we did this
- ▶ Packets coming from the other end:
  - ▶ If this is a response to a modified packet, then rewrite the header and send it to the right private IP address
  - ▶ Otherwise, no change
  - ▶ Ok, wait, how does that work?

# NAT: Network Address Translation

A.k.a. “IP Masquerading”

- ▶ We have a gateway router where
  - ▶ One side is a private network
  - ▶ One side is a “real” address
- ▶ Packets coming from the private network:
  - ▶ Rewrite the packet header
  - ▶ Pretend the packet came from the gateway’s “real” address
  - ▶ Remember that we did this
- ▶ Packets coming from the other end:
  - ▶ If this is a response to a modified packet, then rewrite the header and send it to the right private IP address
  - ▶ Otherwise, no change
  - ▶ Ok, wait, how does that work? **Magic**
    - ▶ Uses TCP or UDP information also
    - ▶ And changes port numbers around


# NAT: Network Address Translation

A.k.a. “IP Masquerading”

- ▶ We have a gateway router where
  - ▶ One side is a private network
  - ▶ One side is a “real” address
- ▶ Packets coming from the private network:
  - ▶ Rewrite the packet header
  - ▶ Pretend the packet came from the gateway’s “real” address
  - ▶ Remember that we did this
- ▶ Packets coming from the other end:
  - ▶ If this is a response to a modified packet, then rewrite the header and send it to the right private IP address
  - ▶ Otherwise, no change
  - ▶ Ok, wait, how does that work? **Magic**
    - ▶ Uses TCP or UDP information also
    - ▶ And changes port numbers around

We will see how to set this up . . . later


## Example Home LAN (2)



Suppose gypsy is a typical **home** router (using NAT)

1. cambot sends packets to 129.186.23.166
2. gypsy rewrites source address as w.x.y.z and forwards to ISP
3. ISP and Internet routers get packets to 129.186.23.166
4. Return packets get to w.x.y.z (that's ISP's job)
5. gypsy rewrites destination address and forwards to cambot

## Example Home LAN (2)



Suppose gypsy is a typical **home** router (using NAT)

1. cambot sends packets to 129.186.23.166
2. gypsy rewrites source address as w.x.y.z and forwards to ISP
3. ISP and Internet routers get packets to 129.186.23.166
4. Return packets get to w.x.y.z (that's ISP's job)
5. gypsy rewrites destination address and forwards to cambot

Fun fact: w.x.y.z might **also be a private address**

# Benefits of NAT

- ▶ Make private networks useful
  - ▶ Can connect to “public” Internet
- ▶ Has **significantly** slowed consumption of IP addresses
  - ▶ **Unique** addresses necessary only for publicly-visible things
- ▶ Give us security and natural firewalling
  - ▶ Would-be crackers cannot attack a private IP address

# Benefits of NAT

- ▶ Make private networks useful
  - ▶ Can connect to “public” Internet
- ▶ Has **significantly** slowed consumption of IP addresses
  - ▶ **Unique** addresses necessary only for publicly-visible things
- ▶ Give us security and natural firewalling
  - ▶ Would-be crackers cannot attack a private IP address
- ▶ Most of Iowa State’s machines use private addresses
  - ▶ Including homework server
  - ▶ That’s why you need to go through the VPN from off campus

## Other reserved addresses

The following blocks of addresses are also reserved:

**0.0.0.0/8** : used for broadcast messages on “current” network

**127.0.0.0/8** : used for “loopback” addresses

▶ Address 127.0.0.1 refers to the local host itself

**255.255.255.255/32** : “limited broadcast” destination address



# IPv5

## Internet Stream Protocol (ST and ST2)

- ▶ First specified in 1979
- ▶ Intended as connection-oriented complement to IPv4
- ▶ Second version (ST2) specified in 1990
- ▶ Final version is ST2+
  - ▶ [RFC 1819](#)
- ▶ ST2 uses packets with “5” in the version field
  - ▶ [RFC 1700](#)
- ▶ Never really caught on
- ▶ Technically, was never known as “IPv5”

# IPv6: The new standard for Internet Protocol

- ▶ Not widely implemented yet
  - ▶ October 2011 estimate:
    - 3% of domain names have IPv6 support
    - 12% of networks have IPv6 support
- ▶ Described in [RFC 2460](#) (December 1998)
- ▶ Main motivation: larger address space
  - ▶ Not the only change, but an important one

# IPv6: The new standard for Internet Protocol

- ▶ Not widely implemented yet
  - ▶ October 2011 estimate:
    - 3% of domain names have IPv6 support
    - 12% of networks have IPv6 support
- ▶ Described in [RFC 2460](#) (December 1998)
- ▶ Main motivation: larger address space
  - ▶ Not the only change, but an important one
- ▶ Why are we still using IPv4? Why so slow to adopt IPv6?

# IPv6: The new standard for Internet Protocol

- ▶ Not widely implemented yet
  - ▶ October 2011 estimate:
    - 3% of domain names have IPv6 support
    - 12% of networks have IPv6 support
- ▶ Described in [RFC 2460](#) (December 1998)
- ▶ Main motivation: larger address space
  - ▶ Not the only change, but an important one
- ▶ Why are we still using IPv4? Why so slow to adopt IPv6?
  - ▶ Because it still works
  - ▶ CIDR and NAT have **greatly extended** the life of IPv4

# IPv6: The new standard for Internet Protocol

- ▶ Not widely implemented yet
  - ▶ October 2011 estimate:
    - 3% of domain names have IPv6 support
    - 12% of networks have IPv6 support
- ▶ Described in [RFC 2460](#) (December 1998)
- ▶ Main motivation: larger address space
  - ▶ Not the only change, but an important one
- ▶ Why are we still using IPv4? Why so slow to adopt IPv6?
  - ▶ Because it still works
  - ▶ CIDR and NAT have **greatly extended** the life of IPv4
  - ▶ Need compatible hardware
  - ▶ Need to update applications (think: Y2K problem)
  - ▶ Routing is different
  - ▶ Not easy to “mix” IPv6 and IPv4
  - ▶ But the end is near for IPv4...

# IPv4 address exhaustion

- ▶ IANA ran out of IPv4 addresses on 31 January 2011
- ▶ Asia's **RIR** (Regional Internet Registry) ran out 15 April 2011
- ▶ Europe's RIR ran out 14 September 2012

Vint Cerf discusses IPv6 in 2014. <https://www.youtube.com/watch?v=17GtmwyvmWE>

# What is in an IPv6 packet header?

- ▶ They are different from IPv4
  - ▶ Larger (40 bytes vs. 20 **or more** bytes)
  - ▶ Simpler (8 fields vs. 13 **or more** fields)
- ▶ Version number (should be 6)
  - ▶ Still first 4 bits of the packet, so there is some hope...
- ▶ Payload length: 16 bits
- ▶ Next Header: 8 bits (replaces “Protocol” field)
  - ▶ Usually specifies the transport layer protocol
- ▶ Hop limit: 8 bits (replaces “Time to Live” field)
- ▶ Source address: **128** bits
- ▶ Destination address: **128** bits
- ▶ New in IPv6: “Flow label” field
- ▶ Missing in IPv6: **checksum** and **fragmentation** fields
  - ▶ IPv6 routers **never** fragment IPv6 packets



# Addresses in IPv6

- ▶ 128 bits long
  - ▶ Number of unique addresses:

# Addresses in IPv6

- ▶ 128 bits long
  - ▶ Number of unique addresses:  $2^{128} \approx$

# Addresses in IPv6

- ▶ 128 bits long
  - ▶ Number of unique addresses:  $2^{128} \approx 3.40 \times 10^{38}$
  - ▶ Recall: IPv4 has  $2^{32} \approx 4.29 \times 10^9$  unique addresses

# Addresses in IPv6

- ▶ 128 bits long
  - ▶ Number of unique addresses:  $2^{128} \approx 3.40 \times 10^{38}$
  - ▶ Recall: IPv4 has  $2^{32} \approx 4.29 \times 10^9$  unique addresses
- ▶ Written as
  - ▶ 8 groups of 16-bit integers (in hex) with “:” separators
  - ▶ Can remove leading zeroes from hex values
  - ▶ Can replace **one** “block of zeroes” with empty string
  - ▶ E.g., the following refer to the same address:
    - ▶ 2001:0db8:0000:0000:0000:ff00:0042:8329
    - ▶ 2001:db8:0:0:0:ff00:42:8329
    - ▶ 2001:db8::ff00:42:8329

# Addresses in IPv6

- ▶ 128 bits long
  - ▶ Number of unique addresses:  $2^{128} \approx 3.40 \times 10^{38}$
  - ▶ Recall: IPv4 has  $2^{32} \approx 4.29 \times 10^9$  unique addresses
- ▶ Written as
  - ▶ 8 groups of 16-bit integers (in hex) with “:” separators
  - ▶ Can remove leading zeroes from hex values
  - ▶ Can replace **one** “block of zeroes” with empty string
  - ▶ E.g., the following refer to the same address:
    - ▶ 2001:0db8:0000:0000:0000:ff00:0042:8329
    - ▶ 2001:db8:0:0:0:ff00:42:8329
    - ▶ 2001:db8::ff00:42:8329
- ▶ Special case: **IPv4-mapped** addresses
  - ▶ First 80 bits are 0
  - ▶ Next 16 bits are 1
  - ▶ Last 32 bits are the IPv4 address
  - ▶ Can write the IPv4 address in the usual **dotted notation**
    - ▶ E.g.: ::ffff:192.168.1.1

# Benefits of IPv6

- ▶ Larger address space (obviously)
- ▶ Simpler packet headers
  - ▶ Easier for routers to process
  - ▶ That translates to: speed
- ▶ NAT becomes unnecessary
  - ▶ Can still have private networks, but:
  - ▶ Private network machines can have a **globally unique** address
  - ▶ This is what a **huge** address space buys you
  - ▶ But there is some controversy about this ...

# Old remote login utilities

- ▶ `rcp`: remote file copy
  - ▶ Copy a file from a remote host to the local one or from the local host to a remote one
- ▶ `rsh`: run a shell, remotely
- ▶ `rlogin`: remote login
- ▶ `telnet`: run a shell, remotely
  - ▶ A little smarter than `rsh` and `rlogin`

# Old remote login utilities

- ▶ `rcp`: remote file copy
  - ▶ Copy a file from a remote host to the local one or from the local host to a remote one
- ▶ `rsh`: run a shell, remotely
- ▶ `rlogin`: remote login
- ▶ `telnet`: run a shell, remotely
  - ▶ A little smarter than `rsh` and `rlogin`
- ▶ **NEVER USE THESE**



# Why is it bad to use rcp, rsh, rlogin, telnet?

# Why is it bad to use rcp, rsh, rlogin, telnet?

- ▶ Usernames and passwords are transmitted in “plain text”
  - ▶ I.e., not encrypted in any way
  - ▶ Anyone watching the network traffic can see
    - ▶ But that's not easy to do, right?
- ▶ Allowed .rlogin configuration files in user's home directory
  - ▶ These can be misused and abused ...
  - ▶ ... potentially allowing **anyone** to login without a password
- ▶ rlogin protocol has no way to authenticate the client
  - ▶ Corrupt clients can fake it and gain access

# Why is it bad to use rcp, rsh, rlogin, telnet?

- ▶ Usernames and passwords are transmitted in “plain text”
  - ▶ I.e., not encrypted in any way
  - ▶ Anyone watching the network traffic can see
    - ▶ But that's not easy to do, right?
- ▶ Allowed .rlogin configuration files in user's home directory
  - ▶ These can be misused and abused ...
  - ▶ ... potentially allowing **anyone** to login without a password
- ▶ rlogin protocol has no way to authenticate the client
  - ▶ Corrupt clients can fake it and gain access
- ▶ scp and ssh are **much** better choices

# How hard is it to watch network traffic?

- ▶ Incredibly easy if you have physical access to the network
- ▶ Use a **packet analyzer**
  - ▶ A.k.a.: *network analyzer, packet sniffer, protocol analyzer*
- ▶ These intercept and log network frames
- ▶ How to intercept frames?
  - ▶ Ethernet cards can run in **promiscuous mode**
    - ▶ Copies **all** frames, not just ones to the card
    - ▶ In shared-media network, we can watch all traffic
  - ▶ Network tap
    - ▶ Hardware device that copies network traffic
    - ▶ ... sometimes on just one link
  - ▶ Wireless ethernet: **just listen**
  - ▶ *None of these* can be detected
- ▶ These all have legitimate uses for network debugging

# Secure replacements

## scp

- ▶ Securely copy files between remote and local hosts

## ssh

- ▶ Securely run a shell on a remote host
- ▶ You can do (secure) X forwarding
  - ▶ Remote X client connects to local X server
  - ▶ Which means the graphics are “forwarded” through ssh


Both of these run on top of **SSL**

- ▶ Secure Sockets Layer
- ▶ Cryptographic protocols on top of IP

# Generating an ssh key

1. Run `ssh-keygen` to generate a key
  - ▶ Can associate a passphrase with the key
  - ▶ Key will have two parts
    - (i) A private part (keep this secret)
    - (ii) A public part
  - ▶ How it works is quite technical, but:
    - ▶ The public key is used to encrypt messages
    - ▶ **Only** the private key can decrypt the message
    - ▶ You cannot determine the private key from its public one
2. Add the **public** key to the appropriate file on the remote host
  - ▶ Usually `~/.ssh/authorized_keys`
3. You should be able to `ssh` into the remote host using the key
  - ▶ You will be prompted for the key's passphrase
  - ▶ You will **not** be prompted for the password on the remote host

An appropriate xkcd comic: <http://xkcd.com/865>



End of lecture