

Network Administration

ComS 252 — Iowa State University

Barry Britt and Andrew Miner

Overview of this lecture

This lecture covers assorted topics in “network administration”

- ▶ Tools to manage a home network
- ▶ Tools to manage a huge network

Overview of this lecture

This lecture covers assorted topics in “network administration”

- ▶ Tools to manage a home network
- ▶ Tools to manage a huge network
- ▶ Large networks can be complicated
 - ▶ So necessarily, these topics are complicated
 - ▶ That is the price of flexibility
- ▶ There is a reason you get these network topics last

What is DHCP?

- ▶ Dynamic Host Configuration Protocol
- ▶ A network protocol for configuring network devices
- ▶ Information that can be sent to a client includes:
 - ▶ An IP address to use
 - ▶ The subnet mask to use
 - ▶ IP addresses for gateway machines
 - ▶ IP addresses for DNS servers
- ▶ Is an open protocol
 - ▶ Described in [RFC 2131](#) (1997) for IPv4
 - ▶ Described in [RFC 3315](#) (2003) for IPv6

Why use DHCP? Why not static IP addresses?

- ▶ A DHCP server is easy to set up
- ▶ DHCP clients are trivial to set up
- ▶ Easy to move clients to other networks
- ▶ Network changes are centralized
- ▶ Fewer IP addresses may be required

Why use DHCP? Why not static IP addresses?

- ▶ A DHCP server is easy to set up
- ▶ DHCP clients are trivial to set up
- ▶ Easy to move clients to other networks
- ▶ Network changes are centralized
- ▶ Fewer IP addresses may be required

But what if I need fixed IP addresses for my server machines?

Why use DHCP? Why not static IP addresses?

- ▶ A DHCP server is easy to set up
- ▶ DHCP clients are trivial to set up
- ▶ Easy to move clients to other networks
- ▶ Network changes are centralized
- ▶ Fewer IP addresses may be required

But what if I need fixed IP addresses for my server machines?

- ▶ DHCP can handle this
 - ▶ Use a static IP for your server machine; or
 - ▶ DHCP can always give the same address to the server machine

How does DHCP work?

- ▶ How can the client contact a server when it knows nothing?

How does DHCP work?

- ▶ How can the client contact a server when it knows nothing?
- ▶ UDP **broadcast** packets are used
 - ▶ Go to all machines on the physical subnet
 - ▶ Servers listen for broadcasts on port 67
 - ▶ Clients listen for broadcasts on port 68

How does DHCP work?

- ▶ How can the client contact a server when it knows nothing?
- ▶ UDP **broadcast** packets are used
 - ▶ Go to all machines on the physical subnet
 - ▶ Servers listen for broadcasts on port 67
 - ▶ Clients listen for broadcasts on port 68
- ▶ Can we use one DHCP server for multiple physical subnets?

How does DHCP work?

- ▶ How can the client contact a server when it knows nothing?
- ▶ UDP **broadcast** packets are used
 - ▶ Go to all machines on the physical subnet
 - ▶ Servers listen for broadcasts on port 67
 - ▶ Clients listen for broadcasts on port 68
- ▶ Can we use one DHCP server for multiple physical subnets?
 - ▶ **Yes**, but will need a **relay server** on each subnet
 - ▶ The relay servers forward messages to the DHCP server

The actual (simplified) DHCP protocol

1. Discovery:

- ▶ Broadcast message sent by client
- ▶ Will include client's MAC address
- ▶ May request to keep its last-used IP address

2. Offer:

- ▶ Broadcast message sent by server, including
 - ▶ IP address
 - ▶ Subnet mask
 - ▶ Server's IP address
 - ▶ How long the address may be used (**lease** duration)
 - ▶ Clients must renew the address before the lease expires

3. Request:

- ▶ Client accepts one offer (may receive several)

4. Acknowledgement:

- ▶ Server acknowledges; IP address is assigned

DHCP server

- ▶ Service `dhcpcd`
 - ▶ Manage as usual with `systemctl`
- ▶ Configuration file: path depends on distribution
 - ▶ `/etc/dhpcpd.conf`
 - ▶ `/etc/dhcp/dhpcpd.conf`
- ▶ Several useful man pages:
 - ▶ `dhcpcd`: overview of the server
 - ▶ `dhcpcd.conf`: configuration file
 - ▶ `dhcp-options`: configuration options
 - ▶ Others!

Format of dhcpd.conf

- ▶ Lines starting with “#” are ignored
- ▶ Files contain statements:
 - ▶ Settings for the server
 - ▶ DHCP options to send to the client
 - ▶ “subnet” block
 - ▶ A block of statements that apply to a subnet
 - ▶ “host” block
 - ▶ A block of statements that apply to a host
 - ▶ “group” block
 - ▶ A way to group things with common settings
 - ▶ There are others
- ▶ Let's look at these...

Settings and options in dhcpd.conf

Useful server setting statements

Specify the default and maximum lease times:

- ▶ `default-lease-time <sec>;`
- ▶ `max-lease-time <sec>;`

Settings and options in dhcpd.conf

Useful server setting statements

Specify the default and maximum lease times:

- ▶ `default-lease-time <sec>;`
- ▶ `max-lease-time <sec>;`

Option statements

- ▶ Have the form `option <optname> <value>;`
- ▶ Will cause the specified information to be sent to the client
- ▶ Some useful option names:
 - `subnet-mask`: Good old subnet mask
 - `broadcast-address`: IP address for broadcast
 - `routers`: IP address of gateway
 - `domain-name-servers`: IP addresses for DNS servers

Statements that apply to subnets or hosts

```
subnet <IP-addr> netmask <mask> { <stmts> }
```

- ▶ May include settings discussed earlier
- ▶ Should include one or more “range” statements:
 - ▶ range <low-IP-addr> <high-IP-addr>;
 - ▶ Specifies a range of addresses to pull from

Statements that apply to subnets or hosts

```
subnet <IP-addr> netmask <mask> { <stmts> }
```

- ▶ May include settings discussed earlier
- ▶ Should include one or more “range” statements:
 - ▶ range <low-IP-addr> <high-IP-addr>;
 - ▶ Specifies a range of addresses to pull from

```
host <name> { <stmts> }
```

- ▶ May include settings discussed earlier
- ▶ Should specify a hardware address:
 - ▶ hardware <HW-type> <HW-addr>;
- ▶ Can specify a fixed address:
 - ▶ fixed-address <IP-addr>;

Example dhcpd.conf for a server controlling 2 subnets

```
# Globals
default-lease-time 600;
max-lease-time 7200;
option subnet-mask 255.255.255.0;

subnet 192.168.123.0 netmask 255.255.255.0 {
    option routers 192.168.123.254;
    range 192.168.123.50 192.168.123.249;
}

subnet 192.168.4.0 netmask 255.255.255.0 {
    option routers 192.168.4.1;
    range 192.168.4.100 192.168.4.249;
}
```

Another example dhcpd.conf

```
# Globals
default-lease-time 3600;
max-lease-time 86400;
option subnet-mask 255.255.255.0;
option routers 192.168.0.1;
option domain-name-servers 192.168.0.1, 192.168.0.2;

subnet 192.168.0.0 netmask 255.255.255.0 {
    range 192.168.0.10 192.168.0.199;
}

host webserver {
    hardware ethernet 00:11:22:33:44:55; # MAC address
    fixed-address 192.168.0.8;
}
```

What is DNS?

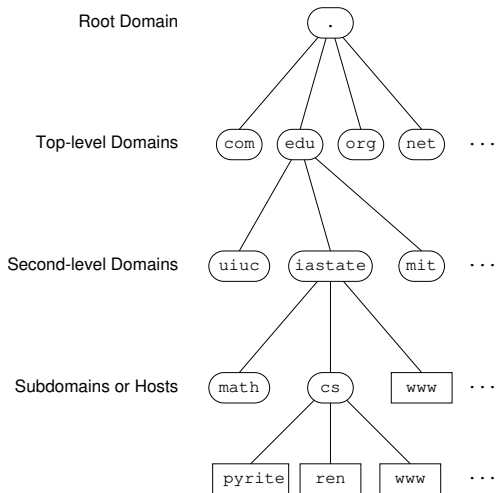
- ▶ Domain Name System
- ▶ IP protocol is based *entirely* on IP addresses
 - ▶ 32 bits for IPv4
 - ▶ 128 bits for IPv6
 - ▶ Packets use these for source and destination addresses
- ▶ But IP addresses are annoying to remember
 - ▶ Will be worse with IPv6
- ▶ Host *names* are easier (e.g., `google.com`, `www.redhat.com`)
- ▶ DNS finds IP addresses for host names
 - ▶ The Internet's "phone book"
- ▶ Is described by **dozens** of RFCs
 - ▶ Basic specs are [RFC 1034](#) and [RFC 1035](#) (1987)

General overview of DNS

- ▶ Conceptually, DNS is a **huge** database
 - ▶ A directory that gives IP addresses for a host name
 - ▶ A *reverse* directory that gives host name(s) for IP address
- ▶ The database is
 - Hierarchical** for quick responses
 - Distributed** for fast access
 - ▶ There are *lots* of DNS servers worldwide
 - ▶ 2011 estimate: 18.5 million DNS servers
 - Replicated** for reliability
- ▶ Today, DNS implementation is
 - Secure** using authentication
 - Fast** using caching

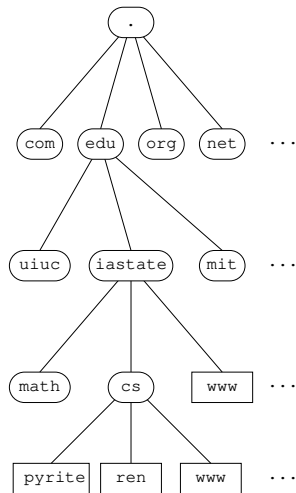
Hierarchical domain structure

- ▶ Each node is a **domain**
- ▶ Top node: **root domain**
- ▶ Node labels: 63 char max
- ▶ A subdomain may have its own subdomains
- ▶ Max. depth: 127
- ▶ **F**ully **Q**ualified **D**omain **N**ame (FQDN): a path from a node to the root
 - ▶ Should include root
 - ▶ E.g., *iastate.edu.*
- ▶ FQDN max length: 255



The Resolver

- ▶ FQDNs are like **absolute paths**
 - ▶ Include the final “.”
- ▶ **Unqualified** domain names:
 - ▶ Do not include the final “.”
 - ▶ Are like **relative paths**
- ▶ The **resolver** (on the client):
 - ▶ Tries to complete unqualified names
 - ▶ Appends domains to produce FQDNs
 - ▶ FQDNs are passed to DNS, one at a time
 - ▶ “.” is always tried first
- ▶ File `/etc/resolv.conf` in Linux:
 - ▶ Configures the resolver
 - ▶ Use **search dom1 dom2 ...**
to specify domains to append, in order
 - ▶ Is usually written by DHCP



Zones

- ▶ A DNS **zone** is a portion of a domain name space
- ▶ Each zone has at least one **authoritative** DNS server
 - ▶ Holds all information about the zone
 - ▶ Knows the DNS servers for any zones “below”
 - ▶ Called “delegation of authority”
 - ▶ Higher-level servers **delegate authority** to lower-level servers
- ▶ There are 13 authoritative DNS servers for the root domain
 - ▶ Operation is regulated by an ICANN committee
 - ▶ Names are A.root-servers.net, ..., M.root-servers.net
 - ▶ Many of these have multiple physical servers
 - ▶ Are distributed worldwide
- ▶ ICANN delegates authority for top-level domains
 - ▶ E.g., there is an authoritative DNS server for edu domain

The three types of DNS queries

1. Iterative queries

- ▶ Client sends a domain name
- ▶ The server does not query other servers
- ▶ The server returns *either*
 - ▶ The IP address for the domain
 - ▶ The name of another DNS server
(either the authoritative one for the domain, or an ancestor)

The three types of DNS queries

1. Iterative queries

- ▶ Client sends a domain name
- ▶ The server does not query other servers
- ▶ The server returns *either*
 - ▶ The IP address for the domain
 - ▶ The name of another DNS server
(either the authoritative one for the domain, or an ancestor)

2. Recursive queries

- ▶ Client sends a domain name
- ▶ The server returns the IP address for the domain
- ▶ The server may query other servers to get this information

The three types of DNS queries

1. Iterative queries

- ▶ Client sends a domain name
- ▶ The server does not query other servers
- ▶ The server returns *either*
 - ▶ The IP address for the domain
 - ▶ The name of another DNS server
(either the authoritative one for the domain, or an ancestor)

2. Recursive queries

- ▶ Client sends a domain name
- ▶ The server returns the IP address for the domain
- ▶ The server may query other servers to get this information

3. Inverse queries

- ▶ Client sends an IP address
- ▶ The server returns a domain name
- ▶ Queries use a special domain named `in-addr.arpa`
 - ▶ E.g., to get the domain name for address 192.168.1.2,
send a request for `2.1.168.192.in-addr.arpa`.

Types of DNS servers

1. Primary (Master) server

- ▶ Is authoritative
- ▶ Holds the “master copy” of the **zone data**
- ▶ Queries answered based on the **zone file** (a local file)
- ▶ Should be **iterative only**
 - ▶ Except maybe for small systems on a trusted private network

Types of DNS servers

1. Primary (Master) server

- ▶ Is authoritative
- ▶ Holds the “master copy” of the **zone data**
- ▶ Queries answered based on the **zone file** (a local file)
- ▶ Should be **iterative only**
 - ▶ Except maybe for small systems on a trusted private network

2. Secondary (Slave) server

- ▶ Is authoritative
- ▶ Obtains zone data from another server (primary or secondary)
- ▶ Can be iterative or recursive

Types of DNS servers

1. Primary (Master) server

- ▶ Is authoritative
- ▶ Holds the “master copy” of the **zone data**
- ▶ Queries answered based on the **zone file** (a local file)
- ▶ Should be **iterative only**
 - ▶ Except maybe for small systems on a trusted private network

2. Secondary (Slave) server

- ▶ Is authoritative
- ▶ Obtains zone data from another server (primary or secondary)
- ▶ Can be iterative or recursive

3. Caching server

- ▶ Is **not** authoritative
- ▶ Store answers to previous queries in cache (memory)
- ▶ Answers requests from cached queries, if it can
- ▶ Otherwise, forwards the query

Resolving a query

- ▶ This is handled transparently by the resolver
- ▶ How does a query to `www.iastate.edu` get resolved?
 1. If the IP address is in the resolver's cache, return it
 2. The resolver contacts the DNS server
 - ▶ Usually, supplied by your ISP
 3. If address is in that server's cache, return it
 4. If this is a **caching** server, forward the query to a “real” server
 5. If address is in that server's cache, return it
 6. When nobody knows, start a query at one of the root servers
 7. Root server sends request to the authoritative server for `edu`.
 8. Request is sent to the authoritative server for `iastate.edu`.
 9. `iastate.edu`. server checks zone file for an entry `www`
- ▶ The **dig** utility can help with debugging DNS queries

Data in a zone file

- ▶ The basic data element in the DNS database is called a **resource record**
- ▶ These are defined in a zone file
- ▶ A resource record contains the following fields:
 - Name** of the node the record refers to
 - TTL** Time to live
 - ▶ This determines how long entries remain cached
 - Class** which should be IN (for Internet)
 - Type** of the record
 - ▶ There are many different types of records
 - ▶ Some of these will be discussed next
 - Data** which varies with the record type

Some types of resource records

A (IPv4 Address)

- Specifies the IPv4 address of a host:

host IN A 192.168.0.200

CNAME (Canonical Name)

- Map an alias to a domain name; e.g.:

ftp IN CNAME www.evilcorp.net.
maps host ftp to www.evilcorp.net.

MX (Mail exchange)

NS (Name server)

- Specify the name server for a domain

PTR (Pointer)

- Used for inverse name resolution, e.g.:

200 IN PTR host

TXT Associates a character string with a domain

One last record type: SOA (Start of Authority)

- ▶ Every zone must have exactly one SOA record
- ▶ Must be the first resource record of the zone file
- ▶ Specifies a nameserver (with an A record)
- ▶ Specifies email address for mail regarding the domain
- ▶ Specifies the following, in order, in parentheses
 - `serial` number
 - ▶ Used to indicate a change in zone data
 - ▶ By convention, this is `yyyymmddnn`:
date of change, and change number
 - `refresh` time (how often to update secondary servers)
 - `retry` time (if a refresh fails, when do we try again)
 - `expire` time (when do secondary servers expire)
 - `minimum` time (how long to cache this zone data)

Simple example zone file

```
; All text after a semi-colon is ignored on a line
$TTL 3D
@      IN      SOA      namesrv.bogus.com.  email.bogus.com.  (
                                2015061101      ; serial number
                                8H              ; refresh time
                                2H              ; retry time
                                4W              ; expire time
                                1D)             ; minimum time
;
      IN      NS       namesrv.bogus.com.
;
namesrv  IN A   192.168.111.1
stan     IN A   192.168.111.15
kyle     IN A   192.168.111.17
cartman  IN A   192.168.111.18
kenny    IN A   192.168.111.22
```

Simple example zone file

```
; All text after a semi-colon is ignored on a line
$TTL 3D
@      IN      SOA      namesrv.bogus.com.  email.bogus.com.  (
                                2015061101      ; serial number
                                8H              ; refresh time
                                2H              ; retry time
                                4W              ; expire time
                                1D)             ; minimum time
;
        IN      NS       namesrv.bogus.com.
;
namesrv  IN A    192.168.111.1
stan     IN A    192.168.111.15
kyle     IN A    192.168.111.17
cartman  IN A    192.168.111.18
kenny    IN A    192.168.111.22
```

Simple example zone file

```
; All text after a semi-colon is ignored on a line
$TTL 3D
@      IN      SOA      namesrv.bogus.com.  email.bogus.com.  (
                                2015061101      ; serial number
                                8H              ; refresh time
                                2H              ; retry time
                                4W              ; expire time
                                1D)             ; minimum time
;
      IN      NS       namesrv.bogus.com.
;
namesrv      IN      A   192.168.111.1
stan         IN      A   192.168.111.15
kyle         IN      A   192.168.111.17
cartman      IN      A   192.168.111.18
kenny        IN      A   192.168.111.22
```

Simple example zone file

```
; All text after a semi-colon is ignored on a line
$TTL 3D
@      IN      SOA      namesrv.bogus.com.  email.bogus.com.  (
                                2015061101      ; serial number
                                8H              ; refresh time
                                2H              ; retry time
                                4W              ; expire time
                                1D)             ; minimum time
;
                                IN      NS      namesrv.bogus.com.
;
namesrv      IN      A      192.168.111.1
stan         IN      A      192.168.111.15
kyle         IN      A      192.168.111.17
cartman      IN      A      192.168.111.18
kenny       IN      A      192.168.111.22
```

Simple example zone file

```
; All text after a semi-colon is ignored on a line
$TTL 3D
@      IN      SOA      namesrv.bogus.com.  email.bogus.com.  (
                                2015061101      ; serial number
                                8H              ; refresh time
                                2H              ; retry time
                                4W              ; expire time
                                1D)             ; minimum time
;
      IN      NS       namesrv.bogus.com.
;
namesrv      IN A      192.168.111.1
stan         IN A      192.168.111.15
kyle         IN A      192.168.111.17
cartman      IN A      192.168.111.18
kenny        IN A      192.168.111.22
```


Simple example zone file

```
; All text after a semi-colon is ignored on a line
$TTL 3D
@      IN      SOA      namesrv.bogus.com.  email.bogus.com.  (
                                2015061101      ; serial number
                                8H              ; refresh time
                                2H              ; retry time
                                4W              ; expire time
                                1D)             ; minimum time
;
      IN      NS       namesrv.bogus.com.
;
namesrv      IN A      192.168.111.1
stan         IN A      192.168.111.15
kyle         IN A      192.168.111.17
cartman      IN A      192.168.111.18
kenny       IN A      192.168.111.22
```

Simple example “inverse” zone file

```
; The inverse zone data for our example domain
$TTL 3D
@      IN      SOA    namesrv.bogus.com.  email.bogus.com.  (
                                2015061101      ; serial number
                                8H              ; refresh time
                                2H              ; retry time
                                4W              ; expire time
                                1D)             ; minimum time
;
      IN      NS      namesrv.bogus.com.
;
1      IN      PTR     namesrv.bogus.com.
15     IN      PTR     stan.bogus.com.
17     IN      PTR     kyle.bogus.com.
18     IN      PTR     cartman.bogus.com.
22     IN      PTR     kenny.bogus.com.
```

Simple example “inverse” zone file

```
; The inverse zone data for our example domain
$TTL 3D
@      IN      SOA    namesrv.bogus.com.  email.bogus.com.  (
                                2015061101      ; serial number
                                8H              ; refresh time
                                2H              ; retry time
                                4W              ; expire time
                                1D)             ; minimum time
;
      IN      NS      namesrv.bogus.com.
;
1      IN      PTR     namesrv.bogus.com.
15     IN      PTR     stan.bogus.com.
17     IN      PTR     kyle.bogus.com.
18     IN      PTR     cartman.bogus.com.
22     IN      PTR     kenny.bogus.com.
```

Setting up a DNS server in Linux

1. Find a decent HOWTO
 - ▶ I like <http://www.tldp.org/HOWTO/DNS-HOWTO.html>
2. Install BIND
 - ▶ Berkeley Internet Name Domain
 - ▶ Is open source (BSD license)
 - ▶ Version 9 is current
3. Edit the `named` configuration file (more on this next...)
 - ▶ In Fedora: `/etc/named.conf`
 - ▶ Make a backup copy first
 - ▶ There are lots of options here; some are important to keep
 - ▶ Zone files will be specified here
4. Create the zone files (if you want an authoritative server)
5. Manage the `named` service as usual

Configuring a DNS server in Linux

Format of `/etc/named.conf`:

- ▶ Text file
- ▶ Sequence of statements of various types
- ▶ Statement syntax:

```
<statement> <arguments> {  
    <option> ;  
    <option> ;  
};
```
- ▶ Statements may be nested
- ▶ Errors in `/etc/named.conf` may prevent `named` from starting
 - ▶ Utility `named-checkconf` can help you here
- ▶ Errors in zone files may prevent `named` from starting
 - ▶ Utility `named-checkzone` can help with these

Configuring a DNS server in Linux (2)

Useful statement types:

`options`

- ▶ Used for global options (at top level)
- ▶ Some useful options:
 - ▶ `listen-on`: Specify subnets where named listens for queries
 - ▶ `allow-query`: Specify hosts or subnets allowed to query
 - ▶ `forwarders`: Specify IP addresses to forward queries to
 - ▶ `forward first`: Ask nameservers listed in `forwarders` first

Configuring a DNS server in Linux (3)

Useful statement types:

`zone zone-name zone-class`

- ▶ Used to define a zone with given name and (optional) class
- ▶ Some useful options:
 - ▶ `file`: absolute or relative pathname for zone file
 - ▶ Relative to `/var/named`
 - ▶ `/etc/named` also a reasonable location
 - ▶ `notify (yes/no)`: Notify secondary servers on update?
 - ▶ `type`: Type for the zone
 - ▶ `master`: This server is authoritative and primary
 - ▶ `slave`: This server is secondary
 - ▶ Others ...

Tiny example named.conf

```
options {  
    listen-on port 53 { 127.0.0.1; 192.168.111.0/24; };  
    :  
    allow-query { localhost; 192.168.111.0/24; };  
    forward first;  
    forwarders { 192.168.50.1; };  
};  
zone "bogus.com" {  
    type master;  
    notify no;  
    file "/etc/named/bogus.com.zone";  
};  
zone "111.168.192.in-addr.arpa" {  
    type master;  
    notify no;  
    file "/etc/named/inverse.bogus.com.zone";  
};
```


Kernel Routing Tables

- ▶ A **Routing Table** is a list of rules for sending packets
- ▶ For each outgoing IP packet, determines **how to send it**
- ▶ The same mechanism for both simple and complex networks
- ▶ Linux supports **multiple** routing tables
 - ▶ Allows routing based on **policy**
 - ▶ Routing table may be selected based on the packet attributes
 - ▶ Uses a **routing policy database**
- ▶ We will discuss the **main** routing table
and assume that all packets are routed the same way

Destination classes

A machine may send packets to

1. Itself

- ▶ Often through a special **loopback** interface
- ▶ These IP addresses are called **local IPs**

2. Locally-connected machines

- ▶ Machines connected to this machine on some LAN
- ▶ These IP addresses are called **directly reachable IPs**

3. Everything else

- ▶ These addresses are reached through one or more gateways

Structure of a routing table

Each routing table entry contains:

- ▶ A destination “subnet”
 - ▶ An IP address
 - ▶ Number of significant bits to match
 - ▶ As an IP mask, usually denoted as Genmask
- ▶ The gateway to use
 - ▶ 0.0.0.0 for none (directly reachable)
 - ▶ Otherwise, IP address of a directly reachable gateway
- ▶ The interface to use

Deciding how to route a packet:

1. Find the rule with the longest matching destination “subnet”
2. Send the packet according to that rule

The kernel does this for **every outgoing IP packet**

Viewing a routing table

netstat utility

- ▶ Shows network stuff
- n : numerical (shows IP addresses)
- r : shows kernel routing tables
- ▶ Is obsolete but still widely used
- ▶ Will display the following flags:
 - U : route is up
 - G : through a gateway
 - ... possibly others

Example routing table

```
prompt$ █
```

Example routing table

```
prompt$ netstat -nr
```

Example routing table

```
prompt$ netstat -nr
Kernel IP routing table
Destination      Gateway          Genmask         Flags ...  Iface
192.168.42.0     0.0.0.0         255.255.255.0   U      ...  eth0
127.0.0.0       0.0.0.0         255.0.0.0       U      ...  lo
0.0.0.0         192.168.42.1   0.0.0.0         UG     ...  eth0
prompt$ █
```

Reading this table:

Example routing table

```
prompt$ netstat -nr
Kernel IP routing table
Destination      Gateway          Genmask         Flags ...  Iface
192.168.42.0     0.0.0.0         255.255.255.0   U      ...  eth0
127.0.0.0       0.0.0.0         255.0.0.0       U      ...  lo
0.0.0.0         192.168.42.1   0.0.0.0         UG     ...  eth0
prompt$
```

Reading this table:

- ▶ Packets for address 192.168.42.x are sent via eth0
 - ▶ This machine must be part of the 192.168.42.0/24 subnet

Example routing table

```
prompt$ netstat -nr
Kernel IP routing table
Destination      Gateway          Genmask          Flags ... Iface
192.168.42.0     0.0.0.0          255.255.255.0    U      ... eth0
127.0.0.0        0.0.0.0          255.0.0.0        U      ... lo
0.0.0.0          192.168.42.1    0.0.0.0          UG     ... eth0
prompt$
```

Reading this table:

- ▶ Packets for address 192.168.42.x are sent via eth0
 - ▶ This machine must be part of the 192.168.42.0/24 subnet
- ▶ Packets for local address 127.x.x.x are sent via lo

Example routing table

```
prompt$ netstat -nr
Kernel IP routing table
Destination      Gateway          Genmask          Flags ... Iface
192.168.42.0     0.0.0.0          255.255.255.0    U      ... eth0
127.0.0.0        0.0.0.0          255.0.0.0        U      ... lo
0.0.0.0          192.168.42.1    0.0.0.0          UG     ... eth0
prompt$
```

Reading this table:

- ▶ Packets for address 192.168.42.x are sent via eth0
 - ▶ This machine must be part of the 192.168.42.0/24 subnet
- ▶ Packets for local address 127.x.x.x are sent via lo
- ▶ Default route: all other packets are sent via eth0
 - ▶ And will go through the gateway 192.168.42.1

Changing a routing table

- ▶ Linux is pretty clever about correctly initializing a routing table
- ▶ If your network interfaces are configured correctly
 - ▶ The routing table should be initialized properly
 - ▶ Usually it is unnecessary to change this

Changing a routing table

- ▶ Linux is pretty clever about correctly initializing a routing table
- ▶ If your network interfaces are configured correctly
 - ▶ The routing table should be initialized properly
 - ▶ Usually it is unnecessary to change this
- ▶ But let's see how to change it, anyway, "by hand"

Changing a routing table

- ▶ Linux is pretty clever about correctly initializing a routing table
- ▶ If your network interfaces are configured correctly
 - ▶ The routing table should be initialized properly
 - ▶ Usually it is unnecessary to change this
- ▶ But let's see how to change it, anyway, “by hand”
- ▶ Old utility: **route**
 - ▶ **route add ...** to add a route
 - ▶ **route del ...** to delete a route
 - ▶ Can also display the routing table

Changing a routing table

- ▶ Linux is pretty clever about correctly initializing a routing table
- ▶ If your network interfaces are configured correctly
 - ▶ The routing table should be initialized properly
 - ▶ Usually it is unnecessary to change this
- ▶ But let's see how to change it, anyway, "by hand"
- ▶ Old utility: **route**
 - ▶ `route add ...` to add a route
 - ▶ `route del ...` to delete a route
 - ▶ Can also display the routing table
- ▶ New utility: **ip**
 - ▶ `ip route add ...` to add a route
 - ▶ `ip route del ...` to delete a route
 - ▶ Can also display the routing table
- ▶ Check your [man](#) pages for details

Properties of a network router

- ▶ A router has more than one network interface
 - ▶ The kernel routing table will be more interesting
 - ▶ Might have different gateways for different destinations
 - ▶ How do we update routing tables in a complex network?
 - ▶ We will discuss this in a minute...

Properties of a network router

- ▶ A router has more than one network interface
 - ▶ The kernel routing table will be more interesting
 - ▶ Might have different gateways for different destinations
 - ▶ How do we update routing tables in a complex network?
 - ▶ We will discuss this in a minute...
- ▶ A router will **forward** packets
 - ▶ From the last example, 192.168.42.1 was the gateway for the 192.168.42.0/24 subnet
 - ▶ Suppose on the gateway machine:
 - ▶ Interface eth0 is on the 192.168.42.0/24 subnet
 - ▶ Interface eth1 is “connected to the Internet”
 - ▶ The gateway must forward packets arriving on eth0 to eth1
 - ▶ We will discuss this in a minute...

Properties of a network router

- ▶ A router has more than one network interface
 - ▶ The kernel routing table will be more interesting
 - ▶ Might have different gateways for different destinations
 - ▶ How do we update routing tables in a complex network?
 - ▶ We will discuss this in a minute...
- ▶ A router will **forward** packets
 - ▶ From the last example, 192.168.42.1 was the gateway for the 192.168.42.0/24 subnet
 - ▶ Suppose on the gateway machine:
 - ▶ Interface eth0 is on the 192.168.42.0/24 subnet
 - ▶ Interface eth1 is "connected to the Internet"
 - ▶ The gateway must forward packets arriving on eth0 to eth1
 - ▶ We will discuss this in a minute...
- ▶ A router may use **NAT** (Network Address Translation)
 - ▶ This is done in the firewall...

Forwarding packets

How to enable packet forwarding in Linux?

Forwarding packets

How to enable packet forwarding in Linux?

- ▶ Remember the `/proc` virtual filesystem?
 - ▶ Dynamically updated, kernel information
- ▶ File `/proc/sys/net/ipv4/ip_forward`
 - ▶ If this contains “1”, the kernel is forwarding packets
 - ▶ If this contains “0”, the kernel is not forwarding packets

Forwarding packets

How to enable packet forwarding in Linux?

- ▶ Remember the `/proc` virtual filesystem?
 - ▶ Dynamically updated, kernel information
- ▶ File `/proc/sys/net/ipv4/ip_forward`
 - ▶ If this contains “1”, the kernel is forwarding packets
 - ▶ If this contains “0”, the kernel is not forwarding packets
- ▶ Extreme magic: we can start or stop forwarding using:
 - ▶ `echo 1 > /proc/sys/net/ipv4/ip_forward` (starts)
 - ▶ `echo 0 > /proc/sys/net/ipv4/ip_forward` (stops)
 - ▶ But this will **not persist** across reboots

Forwarding packets

How to enable packet forwarding in Linux?

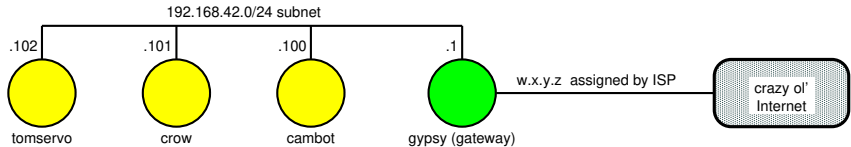
- ▶ Remember the `/proc` virtual filesystem?
 - ▶ Dynamically updated, kernel information
- ▶ File `/proc/sys/net/ipv4/ip_forward`
 - ▶ If this contains “1”, the kernel is forwarding packets
 - ▶ If this contains “0”, the kernel is not forwarding packets
- ▶ Extreme magic: we can start or stop forwarding using:
 - ▶ `echo 1 > /proc/sys/net/ipv4/ip_forward` (starts)
 - ▶ `echo 0 > /proc/sys/net/ipv4/ip_forward` (stops)
 - ▶ But this will **not persist** across reboots
- ▶ To set this value at boot time:
 - ▶ Edit text in file `/etc/sysctl.conf`:
`net.ipv4.ip-forward=1`

Forwarding packets

How to enable packet forwarding in Linux?

- ▶ Remember the `/proc` virtual filesystem?
 - ▶ Dynamically updated, kernel information
- ▶ File `/proc/sys/net/ipv4/ip_forward`
 - ▶ If this contains “1”, the kernel is forwarding packets
 - ▶ If this contains “0”, the kernel is not forwarding packets
- ▶ Extreme magic: we can start or stop forwarding using:
 - ▶ `echo 1 > /proc/sys/net/ipv4/ip_forward` (starts)
 - ▶ `echo 0 > /proc/sys/net/ipv4/ip_forward` (stops)
 - ▶ But this will **not persist** across reboots
- ▶ To set this value at boot time:
 - ▶ Edit text in file `/etc/sysctl.conf`:
`net.ipv4.ip-forward=1`
- ▶ I have also seen these instructions:
 - ▶ Edit text in file `/etc/sysconfig/network`:
`FORWARD_IPV4 = YES`

A LAN gateway



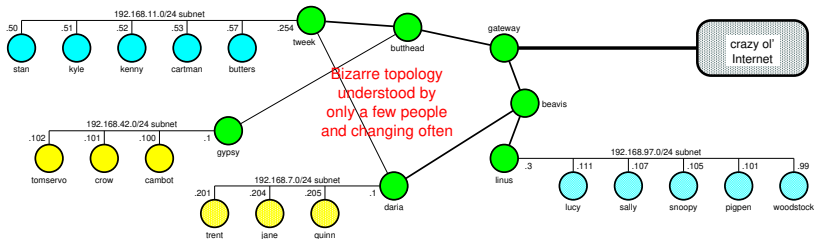
Setting up the kernel routing tables for LAN machines is trivial

- ▶ We already saw an example of this
- ▶ OS usually gets it right with no intervention

Setting up the kernel routing tables for gypsy is trivial

- ▶ Send packets for 192.168.42.0/24 subnet to that interface
- ▶ Send all other packets to Internet interface
- ▶ Use the default gateway given by ISP for the default route
- ▶ OS usually gets this one right with no intervention

What about a collection of LANs?



- ▶ Suppose we have **dozens** or **hundreds** of interconnected LANs
- ▶ E.g., the network on Iowa State campus
- ▶ The topology of routers is complicated
- ▶ Not feasible to set routing tables by hand
- ▶ Not possible for OS to “guess” routing tables
- ▶ Need to do something clever . . .

Routing protocols

- ▶ Routers communicate with each other to determine topology:
 - ▶ Each router knows about networks it is directly attached to
 - ▶ Routers share information with their immediate neighbors
 - ▶ Routers gradually gain knowledge of the network topology
- ▶ Routing algorithms can then choose routes based on topology
- ▶ The kernel routing tables are updated automatically
- ▶ There are different **types** of protocols
 - ▶ Based on the underlying algorithm
 - ▶ For different sizes and types of networks
- ▶ Basic division of protocols:
 1. Interior Gateway Protocols
 2. Exterior Gateway Protocols
- ▶ Before we talk about the Gateway Protocols. . .

Autonomous System (AS)

- ▶ A collection of connected IP routing prefixes (subnets)
- ▶ Under the control of one or more network operators
- ▶ Idea: the individual networks that formed the Internet
- ▶ See [RFC 1930](#) for an in-depth discussion
- ▶ ISU's network is (probably) an Autonomous System

Interior Gateway Protocols (IGPs)

- ▶ An IGP is a routing protocol for use within an AS
- ▶ Two categories based on the algorithm used:
 1. Link-state routing protocol
 2. Distance-vector routing protocol

Link-state routing (an IGP)

- ▶ Each router knows the complete network topology
 - ▶ This is usually learned as part of the protocol
 - ▶ If links fail, the topology is updated

Link-state routing (an IGP)

- ▶ Each router knows the complete network topology
 - ▶ This is usually learned as part of the protocol
 - ▶ If links fail, the topology is updated
- ▶ Each router independently calculates routes
 - ▶ For every possible destination, determine “best next hop”
 - ▶ This is exactly the routing table
 - ▶ “Best next hop” is usually based on lowest latency
 - ▶ Dijkstra’s “shortest path” algorithm will do this

Link–state routing (an IGP)

- ▶ Each router knows the complete network topology
 - ▶ This is usually learned as part of the protocol
 - ▶ If links fail, the topology is updated
- ▶ Each router independently calculates routes
 - ▶ For every possible destination, determine “best next hop”
 - ▶ This is exactly the routing table
 - ▶ “Best next hop” is usually based on lowest latency
 - ▶ Dijkstra’s “shortest path” algorithm will do this
- ▶ Example protocols of this type:
 - ▶ Open Shortest Path First (OSPF)
 - ▶ Common in large enterprise networks
 - ▶ OSPF Version 2 ([RFC 2328](#)) is for IPv4
 - ▶ OSPF Version 3 ([RFC 5340](#)) is for IPv6
 - ▶ Intermediate System to Intermediate System (IS–IS)
 - ▶ Common in large service provider network backbones
 - ▶ Defined by ISO; republished in [RFC 1142](#)

Distance–vector routing protocol (an IGP)

- ▶ Each router does **not** know the complete topology
- ▶ Each router knows its “distance” to immediate neighbors
 - ▶ Distance is either hop count or latency
- ▶ Each router builds a vector of distances to all other routers
- ▶ Each router shares its vector with other routers, periodically
- ▶ Routers update their own vectors based on received ones
 - ▶ This is the **Bellman–Ford** algorithm
 - ▶ The vectors will eventually converge to actual distances
 - ▶ Routing tables are updated based on received vectors

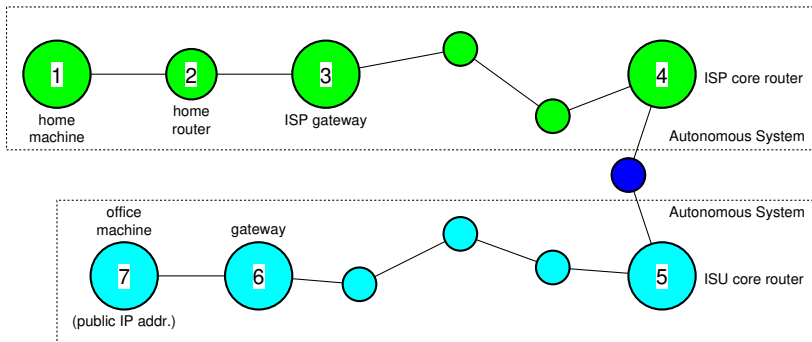
Distance–vector routing protocol (an IGP)

- ▶ Each router does **not** know the complete topology
- ▶ Each router knows its “distance” to immediate neighbors
 - ▶ Distance is either hop count or latency
- ▶ Each router builds a vector of distances to all other routers
- ▶ Each router shares its vector with other routers, periodically
- ▶ Routers update their own vectors based on received ones
 - ▶ This is the **Bellman–Ford** algorithm
 - ▶ The vectors will eventually converge to actual distances
 - ▶ Routing tables are updated based on received vectors
- ▶ Example protocols of this type:
 - ▶ Routing Information Protocol (RIP)
 - ▶ Maximum number of hops allowed is **15**
 - ▶ Version 2 described in [RFC 2453](#) (1998)
 - ▶ Enhanced Interior Gateway Routing Protocol (EIGRP)
 - ▶ Proprietary, by Cisco
 - ▶ Replaces now obsolete IGRP

Exterior Gateway Protocols

- ▶ Each Autonomous System (AS) has one or more **Core Routers**
 - ▶ Low-latency, high-throughput
- ▶ Each AS has an **official** number
- ▶ Core routers use an **Exterior Gateway Protocol** to determine routes between AS's
- ▶ In today's Internet, that protocol is the **Border Gateway Protocol (BGP)**
 - ▶ Described in [RFC 4271](#)

Example: connecting from home to campus



- ▶ Routes 1 → 2 and 2 → 3 are trivial
- ▶ Route 3 → 4 determined by ISP's Interior Gateway Protocol
- ▶ Route 4 → 5 determined by Border Gateway Protocol
- ▶ Route 5 → 6 determined by ISU's IGP
- ▶ Route 6 → 7 is trivial

Firewalld

- ▶ User-space system for dynamic firewall management
- ▶ Discussed basics in earlier lecture
- ▶ This lecture:
 - ▶ A bit more detail how it works
 - ▶ Zones
 - ▶ How to set up NAT

Firewalld architecture

- ▶ Firewalld is essentially a *front-end*
- ▶ Actual packet filtering happens in the kernel
- ▶ Kernel packet filtering framework is [nftables](#)
 - ▶ nftables replaced older iptables in 2014
- ▶ You can set up nftables “by hand”
 - ▶ User-space utility is nft
 - ▶ Gives more control than Firewalld
 - ▶ ...but also more of a pain

Basics of nftables

- ▶ Based on (user-defined) **tables**
- ▶ Each table can have multiple (user-defined) **chains**
- ▶ A chain is a list of **rules**, applied in order
- ▶ A rule typically is
 - ▶ Criteria for matching a packet, such as
 - ▶ Port number
 - ▶ Protocol (e.g., TCP)
 - ▶ Status of a connection
 - ▶ IP address (source or destination)
 - ▶ Combinations of these
 - ▶ A “verdict”: what to do if the packet matches, such as
 - ▶ Accept the packet
 - ▶ Drop the packet
 - ▶ Keep going in the chain
 - ▶ Jump to another chain

Firewalld services

- ▶ Can use `firewall-cmd` to allow or deny *services*
 - ▶ This sets up appropriate rule(s) in appropriate chain and table
 - ▶ Can define your own services
 - ▶ As `.xml` files
 - ▶ Pre-defined usually in `/usr/lib/firewalld/services`
 - ▶ User-defined usually in `/etc/firewalld/services`
- ▶ Can allow/deny services for different **zones**
 - ▶ If not specified, applies to the default zone
- ▶ But what exactly are *zones*?

Firewalld zones

- ▶ A **zone** defines a **level of trust**
 - ▶ Each connection or interface belongs to one zone
 - ▶ A zone may have many connections and interfaces
- ▶ Idea: zones with more trust can allow more services
- ▶ Examples:
 - ▶ Public WiFi connection vs. wired home network connection
 - ▶ Network device connected to “Internet” vs. private network
- ▶ There are pre-defined zones (may vary by distribution)
 - ▶ Typically configured in `/usr/lib/firewalld/zones`
- ▶ You can define your own zones
 - ▶ Typically under `/etc/firewalld/zones`

Some pre-defined zones

- ▶ **drop**: Nothing accepted
 - ▶ All incoming packets dropped; no reply
- ▶ **block**: Nothing accepted
 - ▶ All incoming connections rejected with “icmp-host-prohibited”
- ▶ **public**: Mostly untrusted
 - ▶ Only selected incoming connections are accepted
- ▶ **external**: for routers and NAT
- ▶ **dmz**
- ▶ **work**
- ▶ **home**
- ▶ **internal**: Mostly trusted
 - ▶ Only selected incoming connections are accepted
- ▶ **trusted**: All incoming connections are accepted

Firewall zone configuration

- ▶ Can be done using `firewall-cmd`
 - ▶ `--remove-interface=iface`: remove interface from a zone
 - ▶ `--add-interface=iface`: add interface to a zone
 - ▶ `--get-active-zones`
 - ▶ Show list of active zones
 - ▶ For each zone, show what is using the zone
 - ▶ `--info-zone=zone-name`
 - ▶ Show detailed information for a zone
 - ▶ Includes interfaces and allowed services
- ▶ Can be done “by hand”
 - ▶ Add text “`ZONE=zone-name`” in `ifcfg-iface` file
 - ▶ This assumes the interface always belongs to the zone

NAT with Firewalld

- ▶ Packet forwarding is enabled or disabled as discussed earlier
- ▶ NAT can be turned on for a zone
 - ▶ All forwarded packets that go out on the zone use NAT
 - ▶ Replies are sent back to the right place using NAT
- ▶ `--add-masquerade`: start using NAT on a zone
- ▶ `--remove-masquerade`: stop using NAT on a zone
- ▶ Usually you want to specify a zone with those
(otherwise you get the default zone)
- ▶ Can specify `--permanent` or not as usual
- ▶ `--info-zone=zone-name` shows if masquerading is on/off

Quick example with firewall-cmd

```
prompt$ █
```

Quick example with firewall-cmd

```
prompt$ firewall-cmd --get-active-zones
```

Quick example with firewall-cmd

```
prompt$ firewall-cmd --get-active-zones
home
    interfaces: eth1
external
    interfaces: eth0
prompt$ █
```

Quick example with firewall-cmd

```
prompt$ firewall-cmd --get-active-zones  
home
```

```
    interfaces: eth1  
external
```

```
    interfaces: eth0  
prompt$ firewall-cmd --info-zone=external --permanent
```

Quick example with firewall-cmd

```
prompt$ firewall-cmd --info-zone=external --permanent
external (active)
  target: default
  icmp-block-inversion: no
  interfaces: eth0
  sources:
  services: ssh
  ports:
  protocols:
  masquerade: yes
  forward-ports:
  source-ports:
  icmp-blocks:
  rich rules:

prompt$ █
```

Some utilities

`dig` : DNS lookup utility

`ip` : Update the kernel routing table

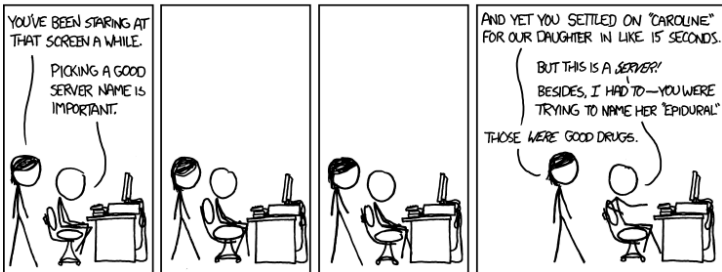
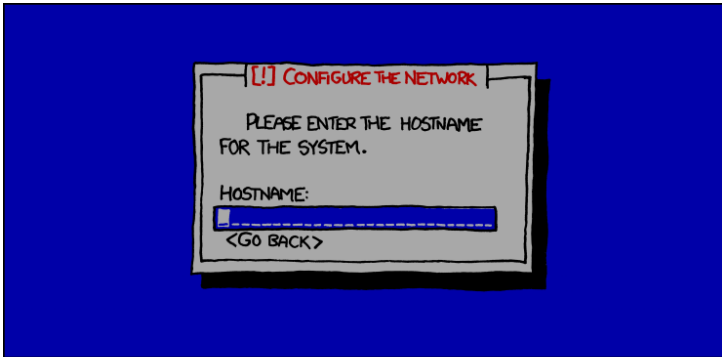
`firewall-cmd` : Update firewall settings

`named-checkconf` : Check named configuration file

`named-checkzone` : Check a named zone file

`netstat` : Show network status

`route` : Update the kernel routing table



End of lecture